

Introduction and Problem Statement

Tanzania, as a developing country, struggles with providing clean water to its population of over 57,000,000. There are many water points already established in the country, but some are in need of repair while others have failed altogether.

The aim of this project is to build a classifier that predicts the condition of a water well (functional, non-functional, or functional but needs repair), using information such as the extraction type, how it is managed, payment type, waterpoint type, the water source, whether it has a permit, and whether a public meeting was held.

With this model, the aim is to help the Government of Tanzania find patterns in non-functional wells to influence how new wells are built.

Objectives

1. Analyze the relationship between the following variables and the `status_group` (functional, non-functional, functional but needs repair) to identify patterns in non-functional wells:
 - `Payment` : What the water costs
 - `source` : The source of the water
 - `management_group` : How the waterpoint is managed
 - `extraction_type` : The kind of extraction the waterpoint uses
 - `permit` : If the waterpoint is permitted
 - `public_meeting` : Whether a public meeting was held for the well
1. Develop a classification model to predict the condition of a well (functional, non-functional, or non-functional but needs repair)

Data Understanding

The original data can be obtained on the [DrivenData 'Pump it Up: Data Mining the Water Table'](#) competition. Basically, there are 4 different data sets; submission format, training set, test set and train labels set which contains status of wells. With given training set and labels set, competitors are expected to build predictive model and apply it to test set to determine status of the wells and submit.

In this project, we will use train set and train label set. Train set has 59400 water points data with 40 features. Train labels data has the same 59400 water points as train set, but just has information about id of these points and status of them.

Limitations of the data

1. The target labels are ternary and highly imbalanced:
 - functional: 32259
 - non functional: 22824

- functional needs repair: 4317

This might impact the performance of the model

1. The data contains 40 features, which are very many, which means that a lot of work will go into EDA and feature elimination

Methods/Data Analysis

After opening the raw training datasets with pandas, we cleaned and prepared the data by imputing missing values, eliminating some redundant columns that contained similar information, identifying the columns with the most variable information for modeling, removing outliers from some numerical columns, encoding categorical variables for modeling.

The 15 final selected features for the model were:

- basin
- month_recorded
- region
- extraction_type_class
- management_group
- payment
- quality_group
- quantity
- source
- waterpoint_type
- gps_height
- population
- construction_age
- permit
- public_meeting

Imports and Data Loading

Imports

In [831...

```
import warnings

# Figures
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import statistics

# Data
import pandas as pd
import numpy as np
from operator import itemgetter
import itertools
from collections import defaultdict
from datetime import datetime
```

```

# Models, metrics, scalers and functionalities
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn import svm

from sklearn.model_selection import cross_val_score, cross_val_predict, cross_validate
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV

from sklearn.metrics import confusion_matrix, accuracy_score

from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.feature_selection import SelectKBest, chi2, mutual_info_classif
from sklearn.feature_selection import RFE, RFECV
from sklearn.feature_selection import SelectFromModel

from sklearn.neighbors import KNeighborsClassifier

from sklearn.decomposition import PCA

from sklearn.tree import DecisionTreeClassifier

from IPython.display import Image
from sklearn.tree import export_graphviz

from sklearn.ensemble import RandomForestClassifier
from scipy.stats import randint

import xgboost
import pandas as pd
import numpy as np

```

Data Loading

```

In [832]: training_set_values = pd.read_csv("data/training-set-values.csv")
training_set_labels = pd.read_csv("data/training-set-labels.csv")
print(training_set_values.shape)
print(training_set_labels.shape)

```

```

(59400, 40)
(59400, 2)

```

Both the training labels and values csvs have 59,400 rows. We will now combine them to deal with one training dataset

```

In [833]: training_data = training_set_values.merge(training_set_labels, on='id')
training_data.head()

```

```

Out[833]:

```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	

5 rows × 41 columns

EDA and Feature Exploration

```
In [834]: training_data.shape
```

Out[834]: (59400, 41)

We now have one dataframe with 59,400 rows and 41 columns, with the 41st column being the status_group column

```
In [835]: training_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                         59400 non-null  object
3   funder                                55765 non-null  object
4   gps_height                            59400 non-null  int64
5   installer                             55745 non-null  object
6   longitude                             59400 non-null  float64
7   latitude                              59400 non-null  float64
8   wpt_name                              59400 non-null  object
9   num_private                           59400 non-null  int64
10  basin                                 59400 non-null  object
11  subvillage                            59029 non-null  object
12  region                                59400 non-null  object
13  region_code                           59400 non-null  int64
14  district_code                         59400 non-null  int64
15  lga                                    59400 non-null  object
16  ward                                  59400 non-null  object
17  population                             59400 non-null  int64
18  public_meeting                        56066 non-null  object
19  recorded_by                           59400 non-null  object
20  scheme_management                     55523 non-null  object
21  scheme_name                           31234 non-null  object
22  permit                                56344 non-null  object
23  construction_year                     59400 non-null  int64
24  extraction_type                       59400 non-null  object
25  extraction_type_group                  59400 non-null  object
26  extraction_type_class                  59400 non-null  object
27  management                             59400 non-null  object
28  management_group                       59400 non-null  object
29  payment                                59400 non-null  object
30  payment_type                           59400 non-null  object
31  water_quality                          59400 non-null  object
32  quality_group                          59400 non-null  object
33  quantity                               59400 non-null  object
34  quantity_group                         59400 non-null  object
35  source                                 59400 non-null  object
36  source_type                           59400 non-null  object
37  source_class                           59400 non-null  object
38  waterpoint_type                       59400 non-null  object
39  waterpoint_type_group                  59400 non-null  object
```

```
40 status_group          59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 19.0+ MB
```

```
In [836]: training_data.columns
```

```
Out[836]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
      'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
      'basin', 'subvillage', 'region', 'region_code', 'district_code', 'lga',
      'ward', 'population', 'public_meeting', 'recorded_by',
      'scheme_management', 'scheme_name', 'permit', 'construction_year',
      'extraction_type', 'extraction_type_group', 'extraction_type_class',
      'management', 'management_group', 'payment', 'payment_type',
      'water_quality', 'quality_group', 'quantity', 'quantity_group',
      'source', 'source_type', 'source_class', 'waterpoint_type',
      'waterpoint_type_group', 'status_group'],
      dtype='object')
```

These are all the columns and their descriptions:

- **amount_tsh** - Total static head (amount water available to waterpoint)
- **date_recorded** - The date the row was entered
- **funder** - Who funded the well
- **gps_height** - Altitude of the well
- **installer** - Organization that installed the well
- **longitude** - GPS coordinate
- **latitude** - GPS coordinate
- **wpt_name** - Name of the waterpoint if there is one
- **num_private** -
- **basin** - Geographic water basin
- **subvillage** - Geographic location
- **region** - Geographic location
- **region_code** - Geographic location (coded)
- **district_code** - Geographic location (coded)
- **lga** - Geographic location
- **ward** - Geographic location
- **population** - Population around the well
- **public_meeting** - True/False
- **recorded_by** - Group entering this row of data
- **scheme_management** - Who operates the waterpoint
- **scheme_name** - Who operates the waterpoint
- **permit** - If the waterpoint is permitted
- **construction_year** - Year the waterpoint was constructed
- **extraction_type** - The kind of extraction the waterpoint uses
- **extraction_type_group** - The kind of extraction the waterpoint uses
- **extraction_type_class** - The kind of extraction the waterpoint uses
- **management** - How the waterpoint is managed
- **management_group** - How the waterpoint is managed
- **payment** - What the water costs
- **payment_type** - What the water costs
- **water_quality** - The quality of the water
- **quality_group** - The quality of the water
- **quantity** - The quantity of water

- **quantity_group** - The quantity of water
- **source** - The source of the water
- **source_type** - The source of the water
- **source_class** - The source of the water
- **waterpoint_type** - The kind of waterpoint
- **waterpoint_type_group** - The kind of waterpoint
- **status_group** - The labels in this dataset with three possible values: functional, non-functional, and functional needs repair

```
In [837... training_data.isna().sum() # to see the null values
```

```
Out[837]:
```

id	0
amount_tsh	0
date_recorded	0
funder	3635
gps_height	0
installer	3655
longitude	0
latitude	0
wpt_name	0
num_private	0
basin	0
subvillage	371
region	0
region_code	0
district_code	0
lga	0
ward	0
population	0
public_meeting	3334
recorded_by	0
scheme_management	3877
scheme_name	28166
permit	3056
construction_year	0
extraction_type	0
extraction_type_group	0
extraction_type_class	0
management	0
management_group	0
payment	0
payment_type	0
water_quality	0
quality_group	0
quantity	0
quantity_group	0
source	0
source_type	0
source_class	0
waterpoint_type	0
waterpoint_type_group	0
status_group	0
dtype: int64	

The columns with missing values are:

Funder : 3635

Installer: 3635

Subvillage: 371

Public meeting: 3334

scheme_management:3877

scheme_name: 28166

permit: 3056

There are some columns which contain null and the same information in the data set.

Now, we will drop one for each because the same values or duplicated values do not affect our target, and when we simplify our data we can run our models easier. These are `management` and `scheme_management`, `payment` and `payment_type`, `quantity` and `quantity_group`, `waterpoint_type` and `waterpoint_type_group`, `extraction_type`, `extraction_type_class`, and `extraction_type_group`, `source`, `source_type`, and `source_class`, `water_quality` and `quality_group`. Let's assess them to confirm.

Dealing with columns that have the same information

Management and Scheme_management

In [838..

```
print(training_data.management.value_counts())
print("-----")
print(training_data.scheme_management.value_counts())
```

```
vwc          40507
wug           6515
water board   2933
wua           2535
private operator 1971
parastatal    1768
water authority  904
other          844
company        685
unknown        561
other - school   99
trust           78
Name: management, dtype: int64
-----
VWC           36793
WUG            5206
Water authority 3153
WUA            2883
Water Board    2748
Parastatal     1680
Private operator 1063
Company         1061
Other           766
SWC             97
Trust           72
None            1
Name: scheme_management, dtype: int64
```

These two columns have nearly the same information. `Scheme_management` represents who operates the water point, `'management'` represents how the water point is managed. There are 3877 null values in `'scheme_management'` column so we prefer to keep the `'management'` column.

In [839..

```
training_data.drop('scheme_management', axis=1, inplace=True)
```

Payment and Payment type

In [840..

```
# Looking at payment and payment_type
print(training_data.payment.value_counts())
```

```
print("-----")
print(training_data.payment_type.value_counts())
```

```
never pay          25348
pay per bucket     8985
pay monthly        8300
unknown            8157
pay when scheme fails 3914
pay annually       3642
other              1054
```

```
Name: payment, dtype: int64
```

```
-----
```

```
never pay      25348
per bucket     8985
monthly        8300
unknown        8157
on failure     3914
annually       3642
other          1054
```

```
Name: payment_type, dtype: int64
```

These two columns have the same exact information, but `payment` has more details on the column naming, so we will keep it and drop `payment_type`

```
In [841... training_data.drop('payment_type', axis=1, inplace=True)
```

Quantity and Quantity_Group

```
In [842... # Looking at `quantity and quantity_group`
print(training_data.quantity.value_counts())
print("-----")
print(training_data.quantity_group.value_counts())
```

```
enough          33186
insufficient     15129
dry              6246
seasonal         4050
unknown          789
Name: quantity, dtype: int64
```

```
-----
```

```
enough          33186
insufficient     15129
dry              6246
seasonal         4050
unknown          789
Name: quantity_group, dtype: int64
```

These have the same exact information, so we can drop either

```
In [843... training_data.drop('quantity_group', axis=1, inplace=True)
```

Waterpoint Type and Waterpoint_type_group

```
In [844... # Looking at the waterpoint_type and waterpoint_type_group columns
print(training_data.waterpoint_type.value_counts())
print("-----")
print(training_data.waterpoint_type_group.value_counts())
```

```
communal standpipe 28522
hand pump          17488
other              6380
communal standpipe multiple 6103
improved spring    784
cattle trough      116
```



```

dam
Name: waterpoint_type, dtype: int64
-----
communal standpipe    34625
hand pump             17488
other                 6380
improved spring       784
cattle trough         116
dam                   7
Name: waterpoint_type_group, dtype: int64

```

`Waterpoint_type` has more information and distribution, so we will keep it and drop `waterpoint_type_group`

```
In [845... training_data.drop(columns = 'waterpoint_type_group', axis=1, inplace=True)
```

Extraction_type, Extraction_type_class and Extraction_type_group columns

```
In [846... # Let's look at the extraction_type, extraction_type_class and extraction_type_group col
print(training_data.extraction_type.value_counts())
```

```

gravity                26780
nira/tanira            8154
other                  6430
submersible            4764
swn 80                 3670
mono                  2865
india mark ii          2400
afridev                1770
ksb                    1415
other - rope pump      451
other - swn 81         229
windmill               117
india mark iii         98
cemo                   90
other - play pump      85
walimi                 48
climax                 32
other - mkulima/shinyanga 2
Name: extraction_type, dtype: int64

```

```
In [847... print(training_data.extraction_type_group.value_counts())
```

```

gravity                26780
nira/tanira            8154
other                  6430
submersible            6179
swn 80                 3670
mono                  2865
india mark ii          2400
afridev                1770
rope pump              451
other handpump         364
other motorpump        122
wind-powered           117
india mark iii         98
Name: extraction_type_group, dtype: int64

```

```
In [848... print(training_data.extraction_type_class.value_counts())
```

```

gravity                26780
handpump              16456
other                 6430
submersible           6179
motorpump             2987

```

```
rope pump          451
wind-powered       117
Name: extraction_type_class, dtype: int64
```

`extraction_type_class` is an even further simplification of `extraction_type`. Since it is the simplest feature of the three, this is the one selected for use, and the handling of category level other will take place after dummification.

```
In [849... training_data.drop('extraction_type_group', axis=1, inplace=True)
```

```
In [850... training_data.drop('extraction_type', axis=1, inplace=True)
```

Source, Source_type_and Source_class

```
In [851... # Let's assess the source, source_type_and source_class columns
# SOURCE
print(training_data.source.value_counts())
```

```
spring          17021
shallow well    16824
machine dbh     11075
river           9612
rainwater harvesting  2295
hand dtw        874
lake            765
dam             656
other           212
unknown         66
Name: source, dtype: int64
```

```
In [852... # SOURCE TYPE
print(training_data.source_type.value_counts())
```

```
spring          17021
shallow well    16824
borehole        11949
river/lake      10377
rainwater harvesting  2295
dam             656
other           278
Name: source_type, dtype: int64
```

```
In [853... # SOURCE CLASS
print(training_data.source_class.value_counts())
```

```
groundwater     45794
surface         13328
unknown         278
Name: source_class, dtype: int64
```

It is obvious that these columns have the same information. We will keep the `source` column and drop the rest since it has more information

```
In [854... training_data.drop('source_type', axis=1, inplace=True)
```

```
In [855... training_data.drop('source_class', axis=1, inplace=True)
```

Water_quality and Quality_group

```
In [856... # Looking at the water_quality and quality_group columns
print(training_data.water_quality.value_counts())
```

```
print("")
print(training_data.quality_group.value_counts())
```

```
soft          50818
salty         4856
unknown       1876
milky         804
coloured      490
salty abandoned 339
fluoride      200
fluoride abandoned 17
Name: water_quality, dtype: int64
-----
```

```
good          50818
salty         5195
unknown       1876
milky         804
colored       490
fluoride      217
Name: quality_group, dtype: int64
```

Since `quality_group` has less categories, we will keep it and drop `water_quality`

```
In [857... training_data.drop('water_quality', axis=1, inplace=True)
```

```
In [858... # Let's look at our new data
training_data.columns
```

```
Out[858]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
        'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
        'basin', 'subvillage', 'region', 'region_code', 'district_code', 'lga',
        'ward', 'population', 'public_meeting', 'recorded_by', 'scheme_name',
        'permit', 'construction_year', 'extraction_type_class', 'management',
        'management_group', 'payment', 'quality_group', 'quantity', 'source',
        'waterpoint_type', 'status_group'],
        dtype='object')
```

```
In [859... training_data.head()
```

```
Out[859]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	

5 rows × 32 columns

Conclusion from dealing with columns that have similar data

We dropped the following columns in favor of those with more robust data:

- `scheme_management`, `payment_type`, `quantity_group`, `waterpoint_type_group`, `extraction_type`, `extration_type_group`, `source_type`, `source_class`,

water_quality

The kept columns are:

- management, payment, quantity, waterpoint, extraction_type_class, source, quality_group

Dealing with Null Values

```
In [860... training_data.isna().sum()
```

```
Out[860]:
```

id	0
amount_tsh	0
date_recorded	0
funder	3635
gps_height	0
installer	3655
longitude	0
latitude	0
wpt_name	0
num_private	0
basin	0
subvillage	371
region	0
region_code	0
district_code	0
lga	0
ward	0
population	0
public_meeting	3334
recorded_by	0
scheme_name	28166
permit	3056
construction_year	0
extraction_type_class	0
management	0
management_group	0
payment	0
quality_group	0
quantity	0
source	0
waterpoint_type	0
status_group	0
dtype: int64	

Let's deal with the columns containing missing values, starting with Funder

Funder

```
In [861... training_data.funder.value_counts()
```

```
Out[861]: Government Of Tanzania    9084  
          Danida                    3114  
          Hesawa                     2202  
          Rwssp                      1374  
          World Bank                 1349  
          ...  
          Kenyans Company            1  
          Rodri                       1  
          Uniceg                      1  
          Tasad                       1
```

```
Mbuzi Mawe 1
Name: funder, Length: 1897, dtype: int64
```

```
In [862... training_data.funder.value_counts()
```

```
Out[862]: Government Of Tanzania    9084
Danida                        3114
Hesawa                        2202
Rwssp                         1374
World Bank                    1349
...
Kenyans Company              1
Rodri                         1
Uniceg                       1
Tasad                        1
Mbuzi Mawe                    1
Name: funder, Length: 1897, dtype: int64
```

The `funder` column has 1897 unique values, hence it will be difficult to fill the null ones. We will drop it

```
In [863... training_data.drop('funder', axis=1, inplace=True)
```

Subvillage

```
In [864... training_data["subvillage"].value_counts()
```

```
Out[864]: Madukani          508
Shuleni        506
Majengo        502
Kati           373
Mtakuja        262
...
Mutoju         1
Migombani B    1
Majaribio      1
Makongoroni    1
Nyuma Ya Mlims 1
Name: subvillage, Length: 19287, dtype: int64
```

```
In [865... len(training_data[training_data["subvillage"].isnull()]["subvillage"])/len(training_data
Out[865]: 0.6245791245791246
```

Given the great scatter (not one subvillage accounts for 1% or more of the examples) this feature is not informative, and therefore dropped.

```
In [866... training_data.drop('subvillage', axis=1, inplace=True)
```

Public Meeting

```
In [867... training_data["public_meeting"].value_counts(dropna = False)/len(training_data)*100
```

```
Out[867]: True      85.877104
False     8.510101
NaN       5.612795
Name: public_meeting, dtype: float64
```

As it is hard to assume whether or not public meetings were actually held, we consider binning this variable into:

```
True
False
```

Unknown

```
In [868... training_data["public_meeting"] = training_data["public_meeting"].astype("category")
training_data["public_meeting"] = training_data["public_meeting"].cat.add_categories('Un
```

```
In [869... training_data.public_meeting.value_counts()
```

```
Out[869]: True      51011
False     5055
Unknown   3334
Name: public_meeting, dtype: int64
```

Scheme Name

```
In [870... training_data["scheme_name"].value_counts(dropna = False)/len(training_data)*100
```

```
Out[870]: NaN      47.417508
K          1.148148
None       1.084175
Borehole   0.919192
Chalinze wate 0.681818
...
NYEHUNGE WATER SUPPLY 0.001684
Kiranjeranje Water supply 0.001684
Kasota 0.001684
BL Motomati 0.001684
Arashi water scheme 0.001684
Name: scheme_name, Length: 2697, dtype: float64
```

`scheme_name` has about half the entries as NaN, and the rest are greatly scattered. Feature `scheme_name` is consequently dropped as well.

```
In [871... training_data.drop('scheme_name', axis=1, inplace=True)
```

Permit

```
In [872... training_data["permit"].value_counts(dropna = False)/len(training_data)*100
```

```
Out[872]: True      65.407407
False     29.447811
NaN        5.144781
Name: permit, dtype: float64
```

`permit` is an intuitive binary feature. About 5% of the values are missing (NaN). It is reasonable to assume that absence of a permit record implies the absence of the permit itself, and hence the NaN values could be imputed as False

```
In [873... training_data = training_data.replace({"permit": {np.nan: False}})
```

```
In [874... training_data.permit.value_counts()
```

```
Out[874]: True      38852
False     20548
Name: permit, dtype: int64
```

Conclusion from dealing with null values

We dropped the following columns:

- `funder`, `subvillage`, and `scheme_name`

The following columns were kept

- `public_meeting` - NaN values were imputed with 'Unknown', introducing a new category
- `permit` - NaN values were imputed with False

Checking for duplicates

```
In [875... training_data.duplicated().sum()
```

```
Out[875]: 0
```

There are no duplicates

Exploring the columns inferred as numeric

```
In [876... # Find numeric variables
numerical_columns = [var for var in training_data.columns if training_data[var].dtype != 'object']

print('There are {} numerical variables\n'.format(len(numerical_columns)))

print('The numerical variables are :', numerical_columns)
```

There are 12 numerical variables

The numerical variables are : ['id', 'amount_tsh', 'gps_height', 'longitude', 'latitude', 'num_private', 'region_code', 'district_code', 'population', 'public_meeting', 'permit', 'construction_year']

```
In [877... training_data[numerical_columns].head(3)
```

```
Out[877]:
```

	id	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
0	69572	6000.0	1390	34.938093	-9.856322	0	11	5	109
1	8776	0.0	1399	34.698766	-2.147466	0	20	2	280
2	34310	25.0	686	37.460664	-3.821329	0	21	4	250

Check for null values in numerical columns

```
In [878... training_data[numerical_columns].isnull().sum()
```

```
Out[878]: id                0
amount_tsh            0
gps_height            0
longitude             0
latitude              0
num_private           0
region_code           0
district_code         0
population            0
public_meeting        0
permit                0
construction_year     0
dtype: int64
```

None of the numeric columns have missing data

```
In [879... training_data[numerical_columns].describe()
```

	id	amount_tsh	gps_height	longitude	latitude	num_private	region_code
count	59400.000000	59400.000000	59400.000000	59400.000000	5.940000e+04	59400.000000	59400.000000
mean	37115.131768	317.650385	668.297239	34.077427	-5.706033e+00	0.474141	15.297003
std	21453.128371	2997.574558	693.116350	6.567432	2.946019e+00	12.236230	17.587406
min	0.000000	0.000000	-90.000000	0.000000	-1.164944e+01	0.000000	1.000000
25%	18519.750000	0.000000	0.000000	33.090347	-8.540621e+00	0.000000	5.000000
50%	37061.500000	0.000000	369.000000	34.908743	-5.021597e+00	0.000000	12.000000
75%	55656.500000	20.000000	1319.250000	37.178387	-3.326156e+00	0.000000	17.000000
max	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	1776.000000	99.000000

Amount tsh

Total static head shows us the height of the flow from the surface. Mostly there are zero values in our dataset. But for zero values no need for pump , because it means we are already in surface.

```
In [880]: training_data["amount_tsh"].value_counts()
```

```
Out[880]: 0.0      41639
          500.0    3102
          50.0    2472
          1000.0   1488
          20.0    1463
          ...
          8500.0     1
          6300.0     1
          220.0      1
          138000.0    1
          12.0       1
          Name: amount_tsh, Length: 98, dtype: int64
```

```
In [881]: training_data["amount_tsh"].value_counts()/len(training_data)*100
```

```
Out[881]: 0.0      70.099327
          500.0     5.222222
          50.0     4.161616
          1000.0    2.505051
          20.0     2.462963
          ...
          8500.0    0.001684
          6300.0    0.001684
          220.0     0.001684
          138000.0  0.001684
          12.0     0.001684
          Name: amount_tsh, Length: 98, dtype: float64
```

We decided to drop this column because 70% of the column has no informative values. So, this column will not be informative to our model and we will drop it.

```
In [882]: training_data.drop('amount_tsh',axis=1, inplace=True)
```

GPS Height

```
In [883]: training_data["gps_height"].value_counts()/len(training_data)*100 # 34% of waterpoint
```

```
Out[883]: 0      34.407407
          -15    0.101010
          -16    0.092593
```



```

-13      0.092593
-20      0.087542
...
2285     0.001684
2424     0.001684
2552     0.001684
2413     0.001684
2385     0.001684
Name: gps_height, Length: 2428, dtype: float64

```

```

In [884... training_data[training_data["gps_height"] > 0]["gps_height"].count()/len(training_data)*
training_data[training_data["gps_height"] < 0]["gps_height"].count()/len(training_data)*

```

```

Out[884]: 2.5185185185185186

```

Gps height shows the level of the water point from sea level. There are 34% zero values but maybe 34% of the water points are at the sea level so we do not change this column now.

Let's use the median

Longitude and Latitude


latitude values make sense (between just below the equator down to 12°S)

longitude values have inconsistencies (0°E is not possible for Tanzania)

```

In [885... training_data.plot(kind='scatter', x="longitude", y="latitude", alpha=0.4, figsize=(14,6)
# To see the outliers

```



```

In [886... training_data.loc[training_data['longitude']==0].head() # to check outside of Tanzania

```

```

Out[886]:

```

	id	date_recorded	gps_height	installer	longitude	latitude	wpt_name	num_private	basin
21	6091	2013-02-10	0	DWE	0.0	-2.000000e-08	Muungano	0	Lake Victoria
53	32376	2011-08-01	0	Government	0.0	-2.000000e-08	Polisi	0	Lake Victoria
168	72678	2013-01-30	0	WVT	0.0	-2.000000e-08	Wvt Tanzania	0	Lake Victoria
177	56725	2013-01-17	0	DWE	0.0	-2.000000e-	Kikundi	0	Lake

						08	Cha Wakina Mama		Victoria
253	13042	2012-10-29	0	DWE	0.0	-2.000000e-08	Kwakisusi	0	Lake Victoria

5 rows × 28 columns

```
In [887]: training_data.loc[training_data['longitude']!=0].describe() # to find the non-zero value
```

```
Out[887]:
```

	id	gps_height	longitude	latitude	num_private	region_code	district_code	
count	57588.00000	57588.000000	57588.000000	57588.000000	57588.000000	57588.000000	57588.000000	575
mean	37106.48807	689.325137	35.149669	-5.885572	0.489060	15.217615	5.728311	1
std	21454.51421	693.564188	2.607428	2.809876	12.426954	17.855254	9.760254	4
min	0.00000	-90.000000	29.607122	-11.649440	0.000000	1.000000	0.000000	
25%	18522.75000	0.000000	33.285100	-8.643841	0.000000	5.000000	2.000000	
50%	37054.50000	426.000000	35.005943	-5.172704	0.000000	12.000000	3.000000	
75%	55667.25000	1332.000000	37.233712	-3.372824	0.000000	17.000000	5.000000	2
max	74247.00000	2770.000000	40.345193	-0.998464	1776.000000	99.000000	80.000000	305

It is obviously seen that it is written as 0 when the longitude is unknown.

Because, the zero points can be seen easily in the graph above outliers and outside of Tanzania.

So, we replaced this with the mean longitude value

```
In [888]: training_data['longitude'].replace(to_replace = 0 , value =35.15, inplace=True) # changi
```

Num Private

num_private column has no description and most of the column contains zeros and is not intuitive to infer.

We will drop it

```
In [889]: training_data["num_private"].value_counts()/len(training_data)*100
```

```
Out[889]:
```

0	98.725589
6	0.136364
1	0.122896
5	0.077441
8	0.077441
...	
180	0.001684
213	0.001684
23	0.001684
55	0.001684
94	0.001684

Name: num_private, Length: 65, dtype: float64

```
In [890]: training_data.drop('num_private',axis=1, inplace=True )
```

Region_code

```
In [891]: training_data["region_code"].nunique()
```

```
Out[891]: 27
```

```
In [892... training_data['region_code'].value_counts()
```

```
Out[892]: 11    5300
          17    5011
          12    4639
           3    4379
           5    4040
          18    3324
          19    3047
           2    3024
          16    2816
          10    2640
           4    2513
           1    2201
          13    2093
          14    1979
          20    1969
          15    1808
           6    1609
          21    1583
          80    1238
          60    1025
          90     917
           7     805
          99     423
           9     390
          24     326
           8     300
          40        1
          Name: region_code, dtype: int64
```

This is a categorical value

District_code

```
In [893... training_data['district_code'].nunique()
```

```
Out[893]: 20
```

```
In [894... training_data['district_code'].value_counts()
```

```
Out[894]: 1    12203
          2    11173
          3     9998
          4     8999
          5     4356
          6     4074
          7     3343
          8     1043
         30      995
         33      874
         53      745
         43      505
         13      391
         23      293
         63      195
         62      109
         60        63
          0        23
         80        12
         67         6
          Name: district_code, dtype: int64
```

This is a categorical value

Population


In the `population` column, 36% of rows have value 0. These could be mistakes or numeric encoding for NA (maybe waterpoints where population ultimately migrated?).

Do we impute these entries? The distribution is considerably skewed.

```
In [895... print(training_data['population'].skew())
```

```
12.660713588843592
```

```
In [896... plt.figure(figsize=(10,4))
training_data.boxplot(['population'], grid=True, fontsize=10)
plt.show()
```



```
In [897... training_data["population"].describe()
```

```
Out[897]: count    59400.000000
mean      179.909983
std       471.482176
min        0.000000
25%        0.000000
50%       25.000000
75%      215.000000
max     30500.000000
Name: population, dtype: float64
```

imputation is preferred to be with the median considering the significant skewness.

```
In [898... population_median = training_data[(training_data["population"] != 0)]["population"].desc
```

```
In [899... training_data["population"] = training_data["population"].replace({0:population_median})
```

```
In [900... training_data.population.value_counts()
```

```
Out[900]: 150      23273
1         7025
200       1940
250       1681
300       1476
...
406         1
1960        1
1685        1
2248        1
```

1439 1
Name: population, Length: 1048, dtype: int64

Permit

```
In [901...] training_data['permit']
```

```
Out[901]: 0      False
          1      True
          2      True
          3      True
          4      True
          ...
          59395   True
          59396   True
          59397   False
          59398   True
          59399   True
          Name: permit, Length: 59400, dtype: bool
```

```
In [902...] training_data["permit"].value_counts(dropna = False)/len(training_data)*100
```

```
Out[902]: True      65.407407
          False    34.592593
          Name: permit, dtype: float64
```

Permit is an intuitive binary feature with Boolean values (True or False)

Construction_year

```
In [903...] print(training_data['construction_year'].skew())
```

-0.6349277865999228

```
In [904...] training_data["construction_year"].value_counts()/len(training_data)*100
```

```
Out[904]: 0      34.863636
          2010    4.452862
          2008    4.398990
          2009    4.264310
          2000    3.520202
          2007    2.671717
          2006    2.476431
          2003    2.164983
          2011    2.114478
          2004    1.890572
          2012    1.824916
          2002    1.809764
          1978    1.745791
          1995    1.707071
          2005    1.702020
          1999    1.648148
          1998    1.626263
          1990    1.606061
          1985    1.590909
          1980    1.365320
          1996    1.365320
          1984    1.311448
          1982    1.252525
          1994    1.242424
          1972    1.191919
          1974    1.138047
          1997    1.084175
          1992    1.077441
          1993    1.023569
```

```

2001    0.909091
1988    0.877104
1983    0.821549
1975    0.735690
1986    0.730640
1976    0.696970
1970    0.691919
1991    0.545455
1989    0.531987
1987    0.508418
1981    0.400673
1977    0.340067
1979    0.323232
1973    0.309764
2013    0.296296
1971    0.244108
1960    0.171717
1967    0.148148
1963    0.143098
1968    0.129630
1969    0.099327
1964    0.067340
1962    0.050505
1961    0.035354
1965    0.031987
1966    0.028620
Name: construction_year, dtype: float64

```

```
In [905]: training_data[(training_data["construction_year"] != 0)]['construction_year'].describe()
```

```

Out[905]: count    38691.000000
mean      1996.814686
std        12.472045
min        1960.000000
25%        1987.000000
50%        2000.000000
75%        2008.000000
max        2013.000000
Name: construction_year, dtype: float64

```

Imputation is preferred to be with the mean of non-zero years, as it is considered that a majority a missing values stem from older decades.

```
In [906]: construction_year_mean = training_data[(training_data["construction_year"] != 0)][ "const
```

```
In [907]: training_data["construction_year"] = training_data["construction_year"].replace({0:const
```

```
In [908]: training_data["construction_year"].value_counts()/len(training_data)*100
```

```

Out[908]: 1996    36.228956
2010     4.452862
2008     4.398990
2009     4.264310
2000     3.520202
2007     2.671717
2006     2.476431
2003     2.164983
2011     2.114478
2004     1.890572
2012     1.824916
2002     1.809764
1978     1.745791
1995     1.707071
2005     1.702020

```

```

1999      1.648148
1998      1.626263
1990      1.606061
1985      1.590909
1980      1.365320
1984      1.311448
1982      1.252525
1994      1.242424
1972      1.191919
1974      1.138047
1997      1.084175
1992      1.077441
1993      1.023569
2001      0.909091
1988      0.877104
1983      0.821549
1975      0.735690
1986      0.730640
1976      0.696970
1970      0.691919
1991      0.545455
1989      0.531987
1987      0.508418
1981      0.400673
1977      0.340067
1979      0.323232
1973      0.309764
2013      0.296296
1971      0.244108
1960      0.171717
1967      0.148148
1963      0.143098
1968      0.129630
1969      0.099327
1964      0.067340
1962      0.050505
1961      0.035354
1965      0.031987
1966      0.028620
Name: construction_year, dtype: float64

```

Conclusion from exploring numeric values

- `Amount_tsh` was dropped
- `GPS Height` remains as it is
- `Longitude` remains as is
- For `latitude`, we imputed the 0 values with the mean since it does not make sense to have a latitude of 0 in Tanzania
- `Num private` was dropped
- `Region code` and `district code` are categorical values and remain as they are
- `Population` needs to be imputed with the median for the years listed as zero
- `Permit` is an intuitive binary feature with Boolean values (True or False)
- `Construction year` will be imputed with the mean of non-zero years for the years listed as zero

Let's look at our new data

```
In [909]: training_data.columns
```

```
Out[909]: Index(['id', 'date_recorded', 'gps_height', 'installer', 'longitude',
        'latitude', 'wpt_name', 'basin', 'region', 'region_code',
        'district_code', 'lga', 'ward', 'population', 'public_meeting',

```

```
'recorded_by', 'permit', 'construction_year', 'extraction_type_class',  
'management', 'management_group', 'payment', 'quality_group',  
'quantity', 'source', 'waterpoint_type', 'status_group'],  
dtype='object')
```

Explore the columns inferred as non-numeric

Initially all these columns are inferred as type object

```
In [910... categorical_columns = [col for col in training_data.columns if training_data[col].dtypes  
categorical_columns
```

```
Out[910]: ['date_recorded',  
            'installer',  
            'wpt_name',  
            'basin',  
            'region',  
            'lga',  
            'ward',  
            'recorded_by',  
            'extraction_type_class',  
            'management',  
            'management_group',  
            'payment',  
            'quality_group',  
            'quantity',  
            'source',  
            'waterpoint_type',  
            'status_group']
```

```
In [911... training_data[categorical_columns].isnull().sum()
```

```
Out[911]: date_recorded          0  
installer          3655  
wpt_name           0  
basin              0  
region             0  
lga                0  
ward               0  
recorded_by        0  
extraction_type_class 0  
management         0  
management_group   0  
payment            0  
quality_group       0  
quantity           0  
source             0  
waterpoint_type     0  
status_group        0  
dtype: int64
```

Date_recorded

We can see that the data type of `Date` variable is object. I will parse the date currently coded as object into datetime format.

```
In [912... training_data["date_recorded"] = pd.to_datetime(training_data["date_recorded"])
```

```
In [913... # extract year from date
```



```
training_data['year_recorded'] = training_data['date_recorded'].dt.year
training_data['year_recorded'].head()
```

```
Out[913]: 0    2011
          1    2013
          2    2013
          3    2013
          4    2011
          Name: year_recorded, dtype: int64
```

```
In [914... training_data['month_recorded'] = training_data['date_recorded'].dt.month

training_data['month_recorded'].head()
```

```
Out[914]: 0     3
          1     3
          2     2
          3     1
          4     7
          Name: month_recorded, dtype: int64
```

```
In [915... # extract day from date

training_data['day_recorded'] = training_data['date_recorded'].dt.day

training_data['day_recorded'].head()
```

```
Out[915]: 0     14
          1      6
          2    25
          3    28
          4    13
          Name: day_recorded, dtype: int64
```

```
In [916... # again view the summary of dataset

training_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    59400 non-null  int64
1   date_recorded         59400 non-null  datetime64[ns]
2   gps_height            59400 non-null  int64
3   installer             55745 non-null  object
4   longitude             59400 non-null  float64
5   latitude              59400 non-null  float64
6   wpt_name              59400 non-null  object
7   basin                59400 non-null  object
8   region                59400 non-null  object
9   region_code           59400 non-null  int64
10  district_code         59400 non-null  int64
11  lga                   59400 non-null  object
12  ward                  59400 non-null  object
13  population            59400 non-null  int64
14  public_meeting        59400 non-null  category
15  recorded_by           59400 non-null  object
16  permit                59400 non-null  bool
17  construction_year     59400 non-null  int64
18  extraction_type_class  59400 non-null  object
19  management            59400 non-null  object
20  management_group      59400 non-null  object
21  payment               59400 non-null  object
22  quality_group         59400 non-null  object
23  quantity              59400 non-null  object
```

```

24 source 59400 non-null object
25 waterpoint_type 59400 non-null object
26 status_group 59400 non-null object
27 year_recorded 59400 non-null int64
28 month_recorded 59400 non-null int64
29 day_recorded 59400 non-null int64
dtypes: bool(1), category(1), datetime64[ns](1), float64(2), int64(9), object(16)
memory usage: 13.3+ MB

```

We can see that there are three additional columns created from `date_recorded` variable. Now, I will drop the original `date_recorded` variable from the dataset.

```
In [917... training_data.drop('date_recorded', axis=1, inplace = True)
```

```
In [918... training_data.head()
```

```
Out[918]:
```

	id	gps_height	installer	longitude	latitude	wpt_name	basin	region	region_code	district_c
0	69572	1390	Roman	34.938093	-9.856322	none	Lake Nyasa	Iringa	11	
1	8776	1399	GRUMETI	34.698766	-2.147466	Zahanati	Lake Victoria	Mara	20	
2	34310	686	World vision	37.460664	-3.821329	Kwa Mahundi	Pangani	Manyara	21	
3	67743	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	Ruvuma / Southern Coast	Mtwara	90	
4	19728	0	Artisan	31.130847	-1.825359	Shuleni	Lake Victoria	Kagera	18	

5 rows × 29 columns

Installer

```
In [919... training_data["installer"].value_counts(dropna = False)/len(training_data)*100
```

```
Out[919]:
```

DWE	29.296296
NaN	6.153199
Government	3.072391
RWE	2.030303
Commu	1.784512
...	
Water Aid/Maji tech	0.001684
Taboma/Community	0.001684
COMPASION INTERNATIO	0.001684
KDPA	0.001684
NZILA	0.001684

Name: installer, Length: 2146, dtype: float64

Given the scatter of installers with only one entry, and that with 2146 different funders it seems complicated to find a common cluster parameter for the specific feature, it is decided to drop installer.

wpt_name

```
In [920... training_data["wpt_name"].value_counts()/len(training_data)*100
```

```
Out[920]:
```

none	5.998316
Shuleni	2.942761
Zahanati	1.397306

```

Msikitini      0.900673
Kanisani       0.543771
...
Kwa Chesi      0.001684
Kwa Padi       0.001684
Kwa Amai       0.001684
Dip Kilulini   0.001684
Kwa Mzee Kunenda 0.001684
Name: wpt_name, Length: 37400, dtype: float64

```

`wpt_name` should play the role of a label but it has null values ("none") and it is not unique (also, the numeric column `id` is already a valid index). Hence, `wpt_name` is dropped.

```
In [921]: training_data.drop('wpt_name', axis=1, inplace = True)
```

Basin

```
In [922]: training_data["basin"].value_counts()/len(training_data)*100
```

```

Out[922]: Lake Victoria      17.252525
Pangani      15.050505
Rufiji       13.427609
Internal     13.106061
Lake Tanganyika 10.828283
Wami / Ruvu  10.079125
Lake Nyasa    8.560606
Ruvuma / Southern Coast 7.563973
Lake Rukwa    4.131313
Name: basin, dtype: float64

```

basin is a nice clustering feature (hydrographic clustering) and has no null values.

Region

```
In [923]: training_data["region"].value_counts()/len(training_data)*100
```

```

Out[923]: Iringa      8.912458
Shinyanga    8.387205
Mbeya        7.809764
Kilimanjaro  7.372054
Morogoro     6.744108
Arusha       5.639731
Kagera       5.582492
Mwanza       5.222222
Kigoma       4.740741
Ruvuma       4.444444
Pwani        4.436027
Tanga        4.287879
Dodoma       3.705387
Singida      3.523569
Mara         3.314815
Tabora       3.297980
Rukwa        3.043771
Mtwara       2.912458
Manyara      2.664983
Lindi        2.602694
Dar es Salaam 1.355219
Name: region, dtype: float64

```

There are 21 distinct values, whereas `region_code` had 27. Either of them could work as geographic clustering feature.

Lga

```
In [924... training_data["lga"].value_counts()/len(training_data)*100
```

```
Out[924]: Njombe          4.213805
          Arusha Rural  2.107744
          Moshi Rural  2.106061
          Bariadi     1.981481
          Rungwe      1.861953
          ...
          Moshi Urban  0.132997
          Kigoma Urban 0.119529
          Arusha Urban 0.106061
          Lindi Urban  0.035354
          Nyamagana    0.001684
          Name: lga, Length: 125, dtype: float64
```

`lga` could be an administrative/political clustering feature, but it has too many levels and classification is not intuitive. Therefore, this feature is dropped.

```
In [925... training_data.drop('lga', axis=1, inplace=True)
```

Ward

```
In [926... training_data["ward"].value_counts()/len(training_data)*100
```

```
Out[926]: Igosi          0.516835
          Imalinyi      0.424242
          Siha Kati     0.390572
          Mdandu        0.388889
          Nduruma       0.365320
          ...
          Mitole        0.001684
          Nyamtinga     0.001684
          Mawenzi       0.001684
          Uchindile     0.001684
          Kinungu       0.001684
          Name: ward, Length: 2092, dtype: float64
```

`ward` seems too scattered to be informative (not one value accounts for even 0.6%) and is therefore dropped.

```
In [927... training_data.drop('ward', axis=1, inplace=True)
```

```
In [928... training_data['public_meeting'].value_counts()
```

```
Out[928]: True          51011
          False         5055
          Unknown       3334
          Name: public_meeting, dtype: int64
```

Recorded by

```
In [929... training_data["recorded_by"].value_counts()/len(training_data)*100
```

```
Out[929]: GeoData Consultants Ltd    100.0
          Name: recorded_by, dtype: float64
```

`recorded_by` comprises only one value, constant across the whole dataset. This feature is uninformative and must be dropped.

```
In [930... training_data.drop('recorded_by', axis=1, inplace=True)
```

Scheme name

Scheme name was dropped when dealing with null values

Extraction type class

This was handled when eliminating redundant categories

Management

This was handled when dealing with redundant categories

```
In [931...] training_data['management'].value_counts()
```

```
Out[931]: vwc          40507
wug          6515
water board  2933
wua          2535
private operator 1971
parastatal   1768
water authority 904
other        844
company      685
unknown      561
other - school 99
trust        78
Name: management, dtype: int64
```

Management Group

```
In [932...] training_data['management_group'].value_counts()
```

```
Out[932]: user-group    52490
commercial    3638
parastatal    1768
other         943
unknown       561
Name: management_group, dtype: int64
```

management_group is chosen over management. It has nulls (0.91%) as unknown that are handled after dummification.

We will drop the management column

```
In [933...] training_data.drop('management', axis=1, inplace=True)
```

Payment

This was dealt with when dealing with redundant columns

Quality Group

This was dealt with when dealing with redundant columns

Quantity

```
In [934...] training_data["quantity"].value_counts()/len(training_data)*100
```

```
Out[934]: enough          55.868687
insufficient    25.469697
dry             10.515152
seasonal        6.818182
```

unknown 1.328283
Name: quantity, dtype: float64

Source

This was dealt with when dealing with redundant columns

Waterpoint Type

This was dealt with when dealing with redundant columns

Conclusion from dealing with non-numeric columns

The following columns were dropped:

- wpt_name , installer , recorded_by , scheme_name

The following columns were kept:

- date_recorded was converted into datetime format
- basin was kept
- region
- lga
- ward
- extraction_type_class
- management_group
- payment
- quality_group
- source
- waterpoint_type

Let's look at our new dataset

In [935... training_data.head()

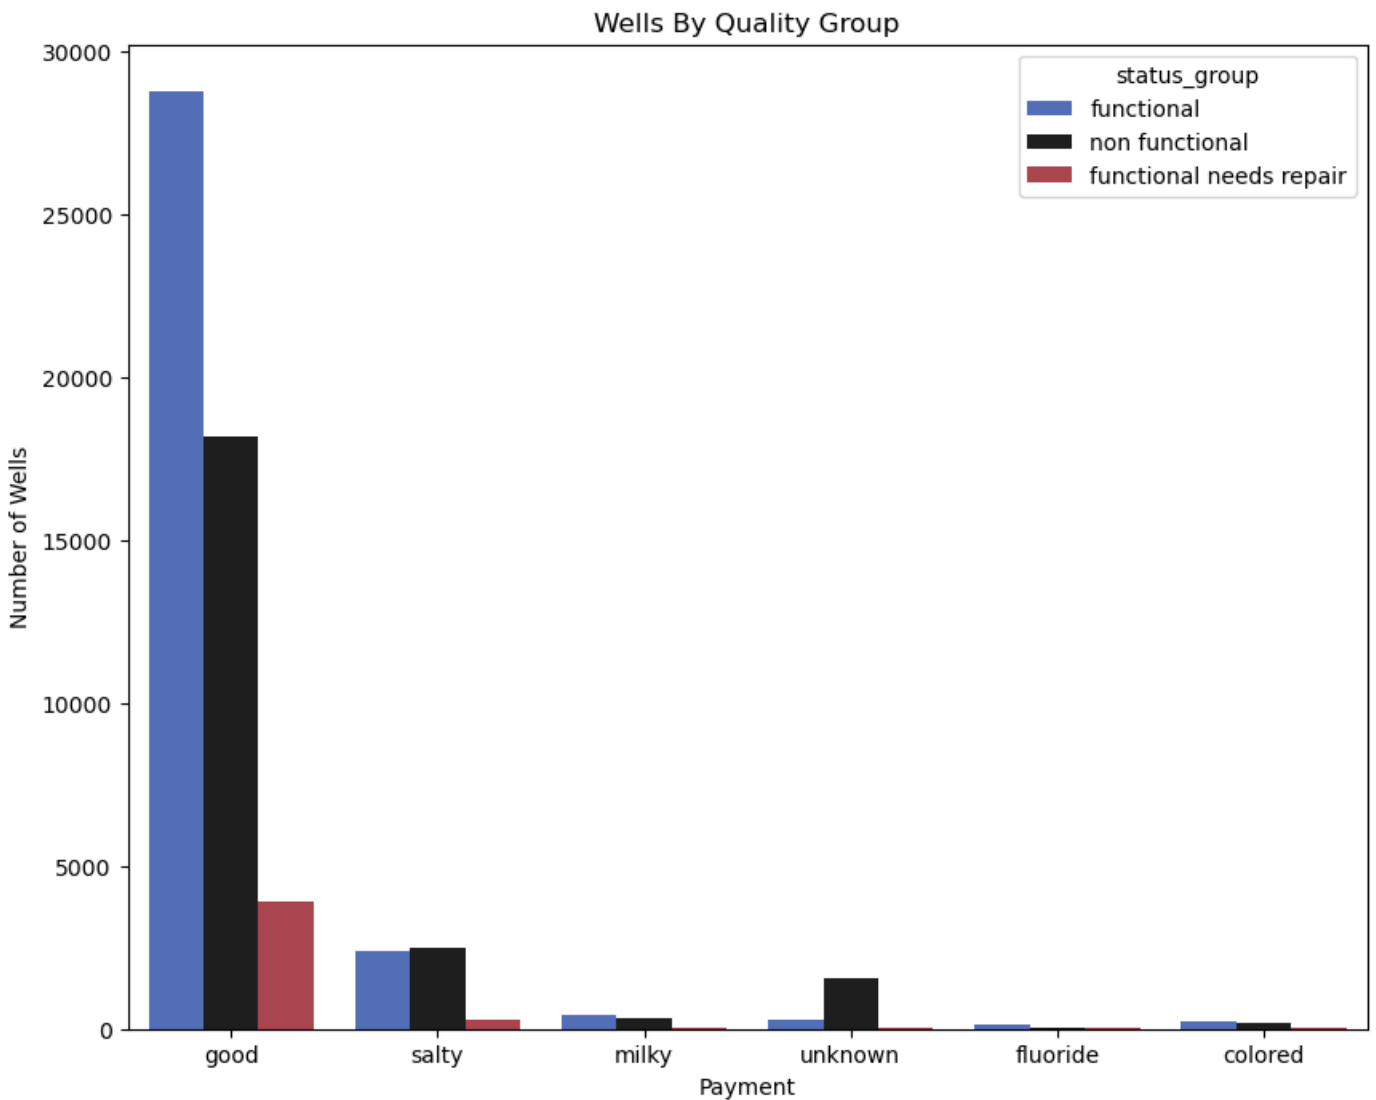
Out[935]:	id	gps_height	installer	longitude	latitude	basin	region	region_code	district_code	popula
0	69572	1390	Roman	34.938093	-9.856322	Lake Nyasa	Iringa	11	5	
1	8776	1399	GRUMETI	34.698766	-2.147466	Lake Victoria	Mara	20	2	
2	34310	686	World vision	37.460664	-3.821329	Pangani	Manyara	21	4	
3	67743	263	UNICEF	38.486161	-11.155298	Ruvuma / Southern Coast	Mtwara	90	63	
4	19728	0	Artisan	31.130847	-1.825359	Lake Victoria	Kagera	18	1	

5 rows × 24 columns

Visualizations

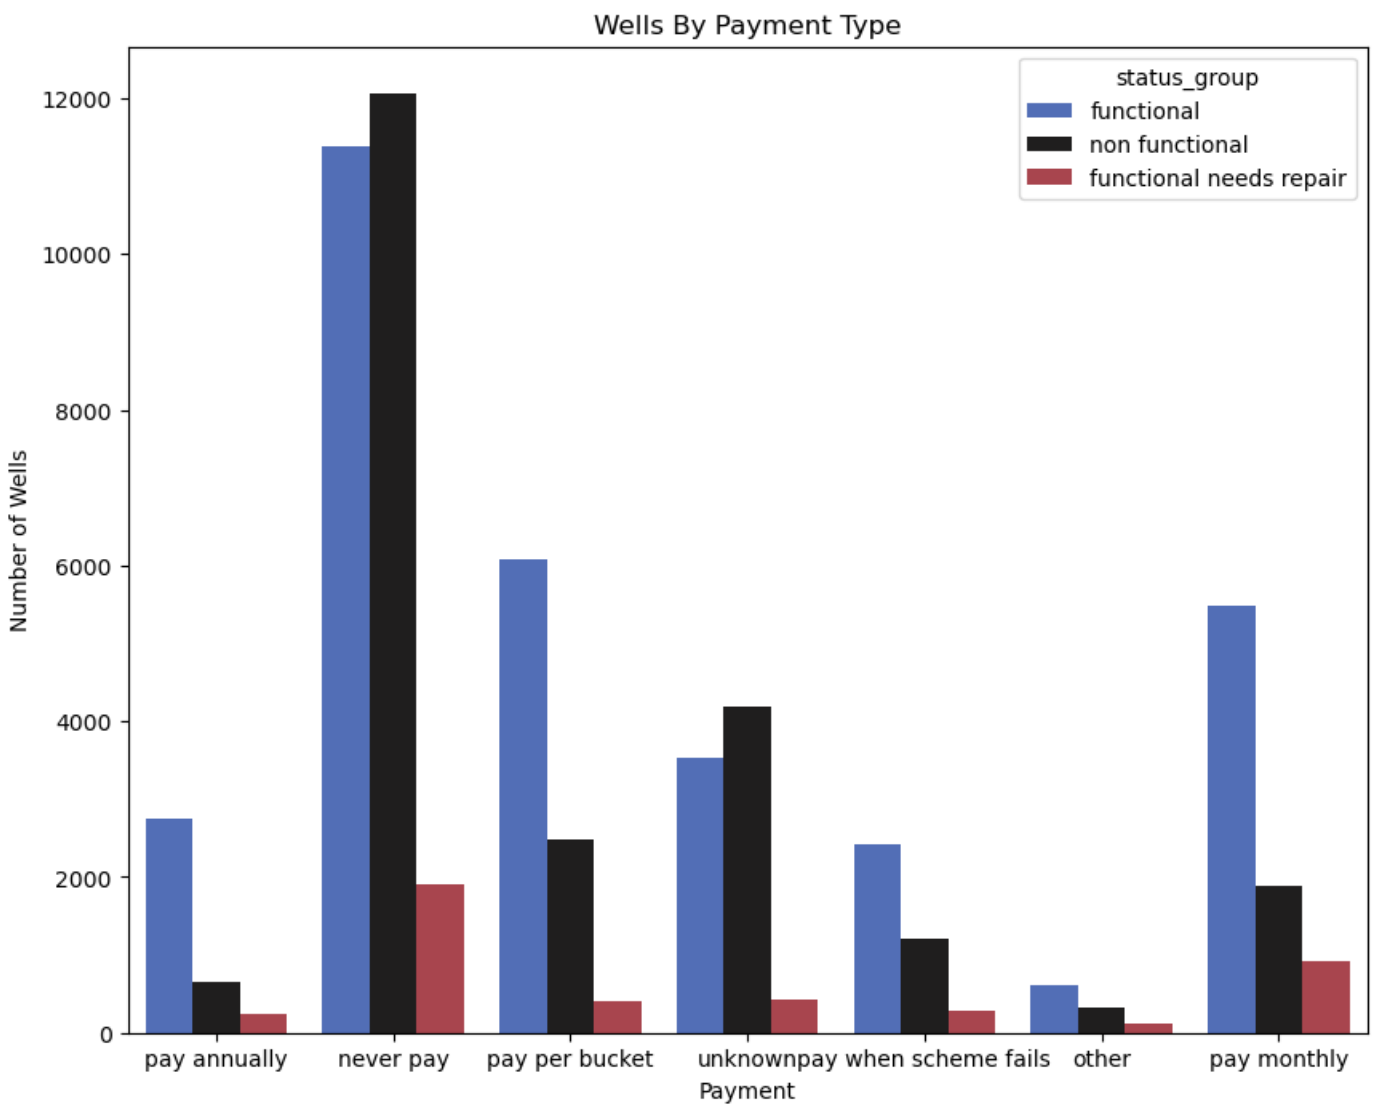
Water Quality by Number of Wells

```
In [936... plt.figure(figsize=(10,8))
ax = sns.countplot(x='quality_group', hue="status_group", data=training_data, palette =
ax.set_xlabel('Payment')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Quality Group');
```




Wells by Payment Type

```
In [937... plt.figure(figsize=(10,8))
ax = sns.countplot(x='payment', hue="status_group", data=training_data, palette = 'icefi
ax.set_xlabel('Payment')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Payment Type');
```




Wells by Waterpoint Type

```
In [938... plt.figure(figsize=(17,10))
ax = sns.countplot(x='waterpoint_type', hue="status_group", data=training_data, palette
ax.set_xlabel('Waterpoint Type')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Waterpoint Type');
```

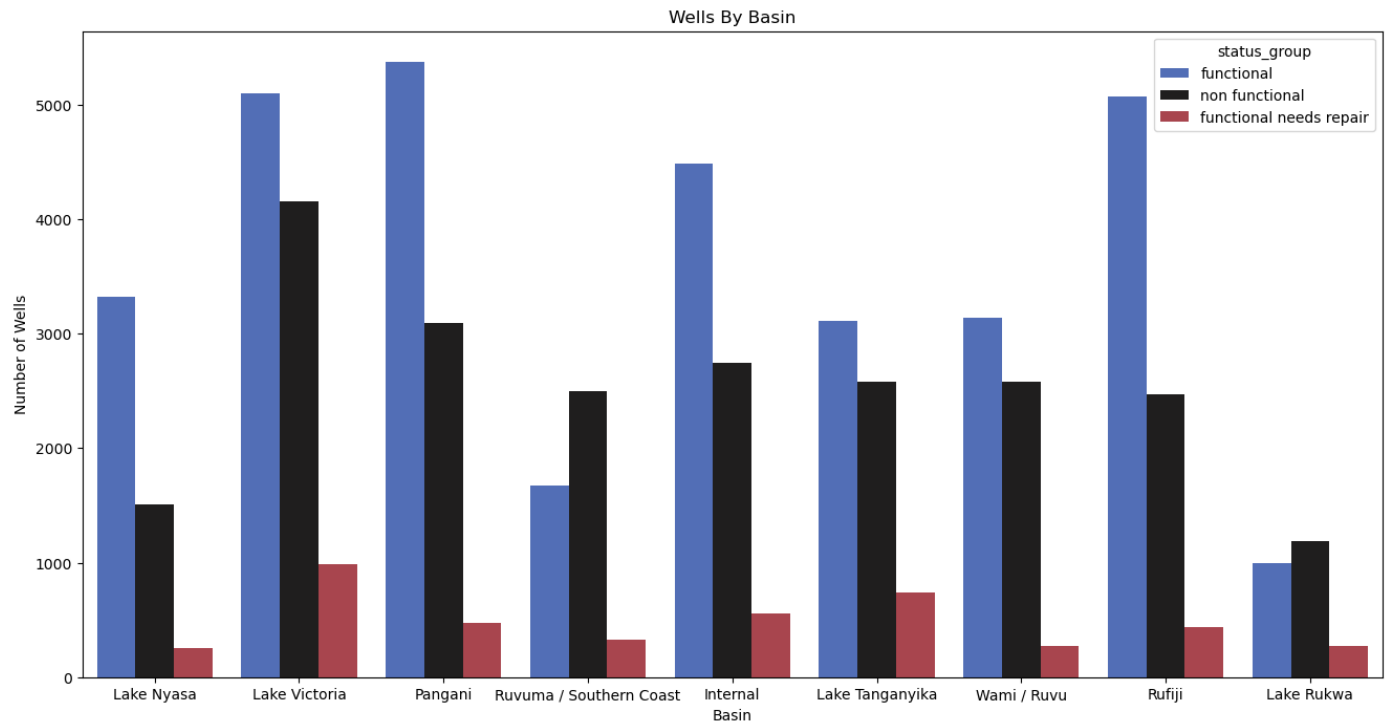
Wells by Quantity

```
In [939... plt.figure(figsize=(10,8))
ax = sns.countplot(x='quantity', hue="status_group", data=training_data, palette = 'icef
ax.set_xlabel('Quantity')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Quantity');
```



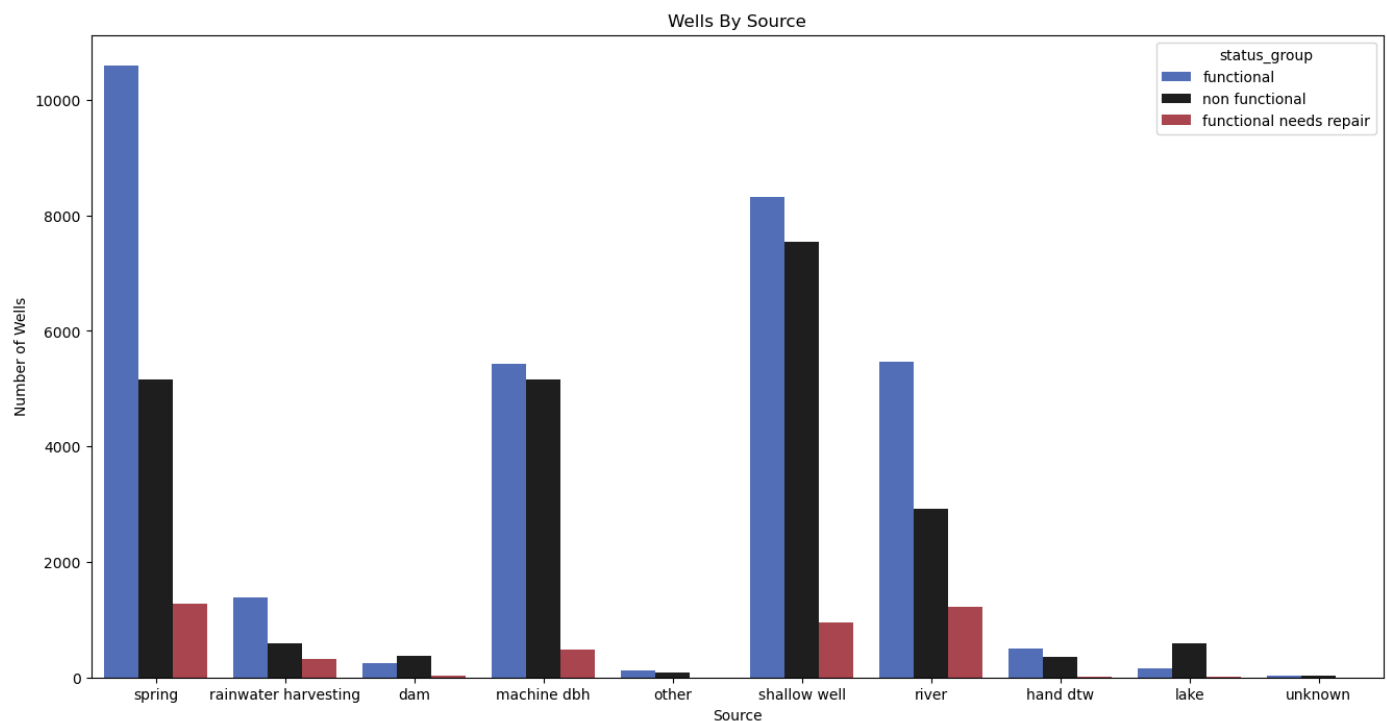
Wells By Basin

```
In [940... plt.figure(figsize=(16,8))
ax = sns.countplot(x='basin', hue="status_group", data=training_data, palette = 'icefire')
ax.set_xlabel('Basin')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Basin');
```



Wells by Source


```
In [941... plt.figure(figsize=(16,8))
ax = sns.countplot(x='source', hue="status_group", data=training_data, palette = 'icefir')
ax.set_xlabel('Source')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Source');
```



Wells By Management Group

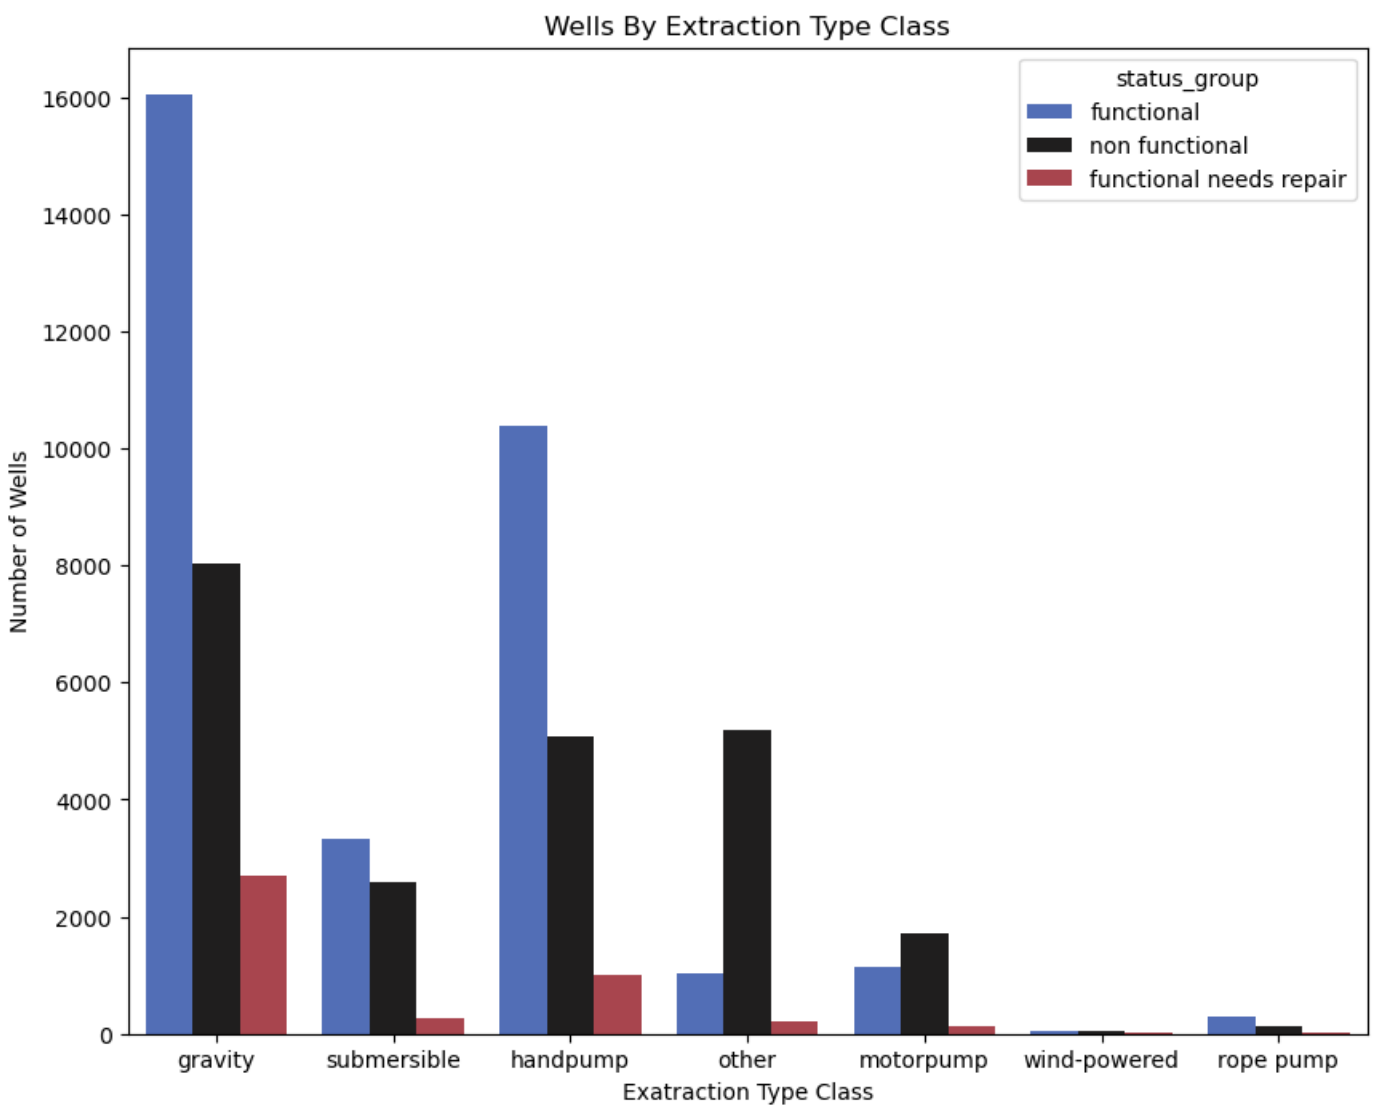
```
In [942... plt.figure(figsize=(10,8))
ax = sns.countplot(x='management_group', hue="status_group", data=training_data, palette = 'icefir')
ax.set_xlabel('Management Group')
```

```
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Management Group');
```




Wells By Extraction Type Class

```
In [943... plt.figure(figsize=(10,8))
ax = sns.countplot(x='extraction_type_class', hue="status_group", data=training_data, pa
ax.set_xlabel('Exatraction Type Class')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Extraction Type Class');
```



Wells By Permit

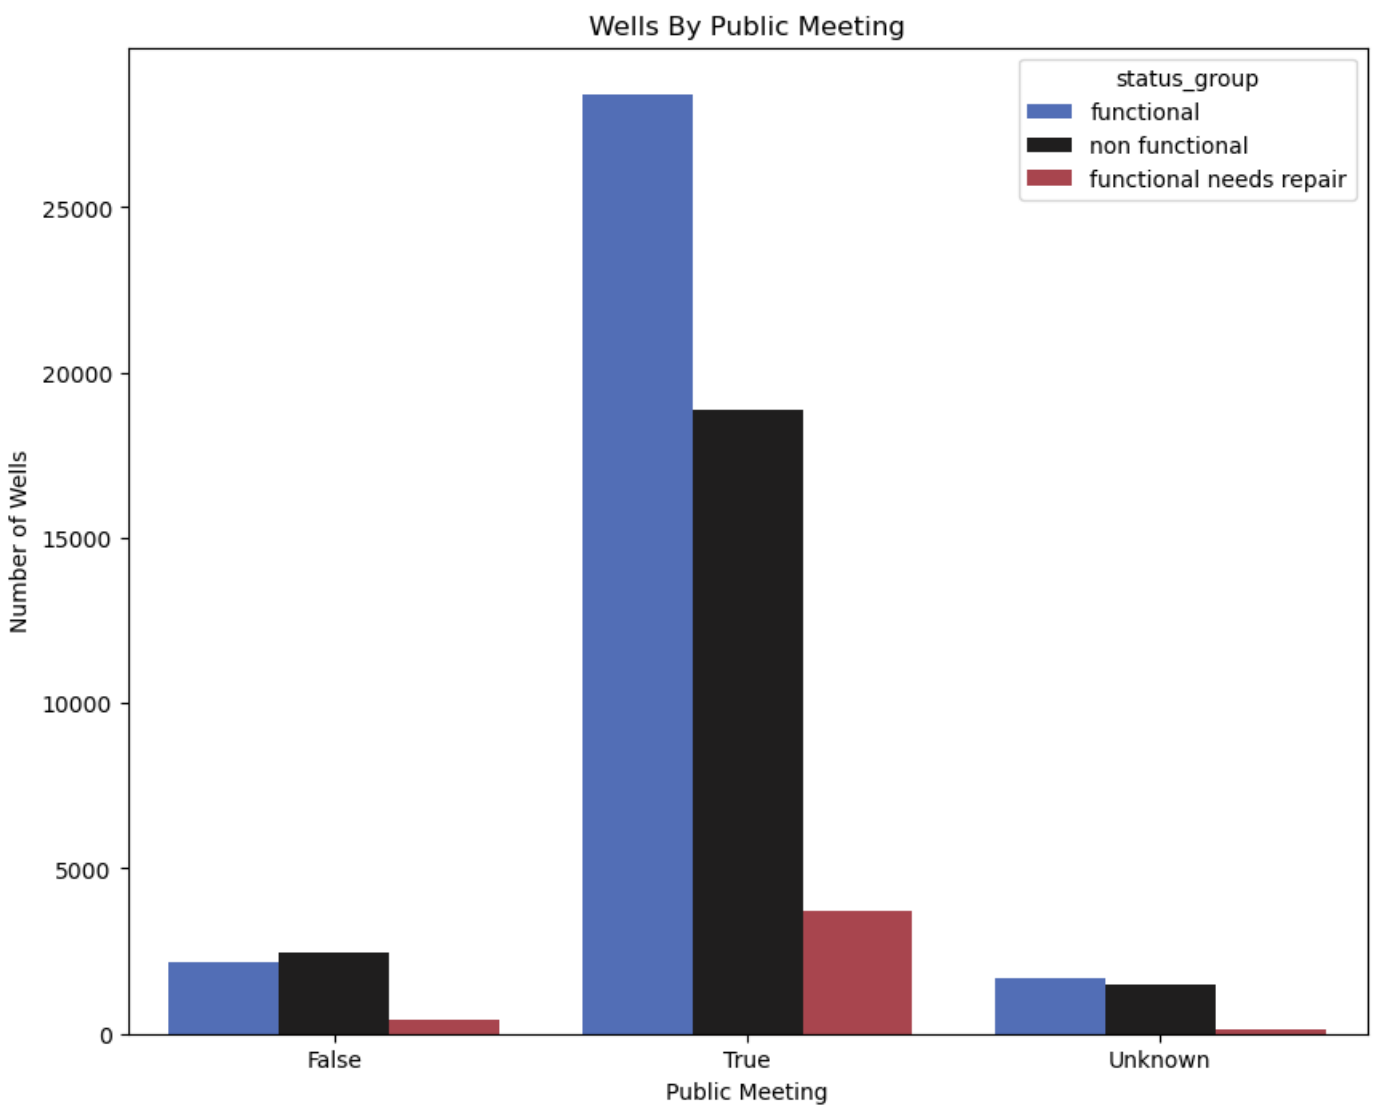
```
In [944... plt.figure(figsize=(10,8))
ax = sns.countplot(x='permit', hue="status_group", data=training_data, palette = 'icefir
ax.set_xlabel('Permit')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Permit');
```



Wells By Public Meeting

In [945...

```
plt.figure(figsize=(10,8))
ax = sns.countplot(x='public_meeting', hue="status_group", data=training_data, palette =
ax.set_xlabel('Public Meeting')
ax.set_ylabel('Number of Wells')
ax.set_title('Wells By Public Meeting');
```



Modeling

Create an initial dataframe dropping all the uninformative features, as well as the categorical features with a large number of levels:

```
In [946...] initial_numeric_cols = ["id", "gps_height",  
                                "population",  
                                "construction_year"]
```

```
In [947...] initial_object_cols = ["basin",  
                                    "month_recorded",  
                                    "region",  
                                    "public_meeting",  
                                    "permit",  
                                    "extraction_type_class",  
                                    "management_group",  
                                    "payment",  
                                    "quality_group",  
                                    "quantity",  
                                    "source",  
                                    "waterpoint_type"]
```

```
In [948...] Xy_init = training_data[initial_numeric_cols + initial_object_cols + ["status_group"]].c
```

```
In [949...] Xy_init = Xy_init.set_index("id")
```

Check missing values

```
In [950... Xy_init.isnull().sum()
```

```
Out[950]: gps_height      0
population    0
construction_year  0
basin         0
month_recorded  0
region        0
public_meeting  0
permit        0
extraction_type_class  0
management_group  0
payment       0
quality_group  0
quantity      0
source        0
waterpoint_type  0
status_group   0
dtype: int64
```

No missing values

```
In [951... Xy_init.head()
```

```
Out[951]:
```

	gps_height	population	construction_year	basin	month_recorded	region	public_meeting	permit
id								
69572	1390	109	1999	Lake Nyasa	3	Iringa	True	False
8776	1399	280	2010	Lake Victoria	3	Mara	Unknown	True
34310	686	250	2009	Pangani	2	Manyara	True	True
67743	263	58	1986	Ruvuma / Southern Coast	1	Mtwara	True	True
19728	0	150	1996	Lake Victoria	7	Kagera	True	True

Feature Engineering

Outliers in numeric variables

```
In [952... training_data[initial_numeric_cols].isnull().sum()
```

```
Out[952]: id      0
gps_height  0
population  0
construction_year  0
dtype: int64
```

For `construction_year`, it was decided to impute the zero years with the mean. Also, it may be more suitable to derive a numeric feature `construction_age` expressing the time delta since construction instead of the original feature.


```
In [953... Xy_init.construction_year.isnull().sum()
```

```
Out[953]: 0
```

```
In [954... current = datetime.now()
```

```
In [955... Xy_init["construction_age"] = (current.year - Xy_init["construction_year"]).astype("int"
```

```
In [956... Xy_init.construction_age.value_counts()
```

```
Out[956]:
```

28	21520
14	2645
16	2613
15	2533
24	2091
17	1587
18	1471
21	1286
13	1256
20	1123
12	1084
22	1075
46	1037
29	1014
19	1011
25	979
26	966
34	954
39	945
44	811
40	779
42	744
30	738
52	708
50	676
27	644
32	640
31	608
23	540
36	521
41	488
49	437
38	434
48	414
54	411
33	324
35	316
37	302
43	238
47	202
45	192
51	184
11	176
53	145
64	102
57	88
61	85
56	77
55	59
60	40
62	30
63	21
59	19

```
58         17
Name: construction_age, dtype: int64
```

Drop `construction_year`

```
In [957... Xy_init.drop('construction_year', axis=1, inplace=True)
```

```
In [958... clean_numeric_cols = [
                                "gps_height",
                                "population",
                                "construction_age"]
clean_numeric_cols
```

```
Out[958]: ['gps_height', 'population', 'construction_age']
```

```
In [959... print(round(Xy_init[clean_numeric_cols].describe()),2)
```

	gps_height	population	construction_age
count	59400.0	59400.0	59400.0
mean	668.0	234.0	27.0
std	693.0	456.0	10.0
min	-90.0	1.0	11.0
25%	0.0	100.0	20.0
50%	369.0	150.0	28.0
75%	1319.0	215.0	28.0
max	2770.0	30500.0	64.0

```
In [960... Xy_init.construction_age.isnull().sum()
```

```
Out[960]: 0
```



I will draw boxplots to visualise outliers in the above variables.

```
In [961... plt.figure(figsize=(20,8))

plt.subplot(2, 2, 1)
fig = Xy_init.boxplot(column='gps_height')
fig.set_title('')
fig.set_ylabel('GPS Height')

plt.subplot(2, 2, 2)
fig = Xy_init.boxplot(column='population')
fig.set_title('')
fig.set_ylabel('Population')

plt.subplot(2, 2, 3)
fig = Xy_init.boxplot(column='construction_age')
fig.set_title('')
fig.set_ylabel('Construction Age');
```



Check the distribution of variables

Now, I will plot the histograms to check distributions to find out if they are normal or skewed. If the variable follows normal distribution, then I will do **Extreme Value Analysis** otherwise if they are skewed, I will find IQR (Interquartile range).




```
In [962... # plot histogram to check distribution

plt.figure(figsize=(15,10))

plt.subplot(2, 2, 1)
fig = Xy_init["gps_height"].hist(bins=10)
fig.set_xlabel('GPS Height')
fig.set_ylabel('Status Group')

plt.subplot(2, 2, 2)
fig = Xy_init["population"].hist(bins=10)
fig.set_xlabel('Population')
fig.set_ylabel('Status Group')

plt.subplot(2, 2, 3)
fig = Xy_init["construction_age"].hist(bins=10)
fig.set_xlabel('Construction Age')
fig.set_ylabel('Status Group');
```



We can see that all the four variables are skewed. So, I will use interquartile range to find outliers.

GPS Height

```
In [963... # find outliers for GPS Height variable

IQR = Xy_init["gps_height"].quantile(0.75) - Xy_init.gps_height.quantile(0.25)
Lower_fence = Xy_init["gps_height"].quantile(0.25) - (IQR * 3)
Upper_fence = Xy_init["gps_height"].quantile(0.75) + (IQR * 3)
print('GPS Height outliers are values < {lowerboundary} or > {upperboundary}'.format(low
GPS Height outliers are values < -3957.75 or > 5277.0
```

For `gps_height`, the minimum and maximum values are -90.0 and 2770.0. So, the outliers are values > 5277. There are no outliers

Population

```
In [964... # find outliers for Population variable

IQR = Xy_init["population"].quantile(0.75) - Xy_init.population.quantile(0.25)
Lower_fence = Xy_init["population"].quantile(0.25) - (IQR * 3)
Upper_fence = Xy_init["population"].quantile(0.75) + (IQR * 3)
print('population outliers are values < {lowerboundary} or > {upperboundary}'.format(low
population outliers are values < -245.0 or > 560.0
```

For `population`, the minimum and maximum values are 1.0 and 30500.0. So, the outliers are values > 560.

Construction Age

```
In [965... # find outliers for Construction Age Variable
```

```

IQR = Xy_init["construction_age"].quantile(0.75) - Xy_init.construction_age.quantile(0.25)
Lower_fence = Xy_init["construction_age"].quantile(0.25) - (IQR * 3)
Upper_fence = Xy_init["construction_age"].quantile(0.75) + (IQR * 3)
print('Construction Age outliers are values < {lowerboundary} or > {upperboundary}'.format(

```

Construction Age outliers are values < -4.0 or > 52.0

For `construction_age`, the minimum and maximum values are 11.0 and 64.0. So, the outliers are values > 52.0.

```

In [966... Xy_init['permit'] = Xy_init['permit'].astype(bool).astype(int) #converting from T/F to 0

```

```

In [967... Xy_init['public_meeting'] = Xy_init['public_meeting'].astype(bool).astype(int) #converting

```

Re-declaring numeric and categorical columns. Moving `permit` and `public_meeting` to numeric now that they are 0-1

```

In [968... initial_numeric_cols = ["id",
                                "gps_height",
                                "population",
                                "construction_year",
                                'permit',
                                'public_meeting']

```

```

In [969... initial_object_cols = ['basin',
                                'month_recorded',
                                'region',
                                'extraction_type_class',
                                'management_group',
                                'payment',
                                'quality_group',
                                'quantity',
                                'source',
                                'waterpoint_type']

```

```

In [970... Xy_init[initial_object_cols] = Xy_init[initial_object_cols].astype("category")

```

```

In [971... Xy_init.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 69572 to 26348
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gps_height             59400 non-null  int64
1   population             59400 non-null  int64
2   basin                  59400 non-null  category
3   month_recorded         59400 non-null  category
4   region                 59400 non-null  category
5   public_meeting         59400 non-null  int32
6   permit                 59400 non-null  int32
7   extraction_type_class  59400 non-null  category
8   management_group       59400 non-null  category
9   payment                59400 non-null  category
10  quality_group           59400 non-null  category
11  quantity                59400 non-null  category
12  source                  59400 non-null  category
13  waterpoint_type        59400 non-null  category
14  status_group            59400 non-null  object
15  construction_age        59400 non-null  int32

```

```
dtypes: category(10), int32(3), int64(2), object(1)
memory usage: 3.1+ MB
```

Declare feature vector and target variable

```
In [972... target_status_group = {'functional':0,
                          'non functional': 2,
                          'functional needs repair': 1}
Xy_init['status_group'] = Xy_init['status_group'].replace(target_status_group)
```

To make our model, we changed the target variable to 0,1 and 2 values.

```
In [973... training_data['status_group'].value_counts()
```

```
Out[973]: functional          32259
non functional          22824
functional needs repair   4317
Name: status_group, dtype: int64
```

```
In [974... X = Xy_init.drop(['status_group'], axis=1)
y = Xy_init['status_group']
```

9. Split data into separate training and test set

```
In [975... # split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state
```

```
In [976... # check the shape of X_train and X_test

X_train.shape, X_test.shape
```

```
Out[976]: ((47520, 15), (11880, 15))
```

```
In [977... # Check data types in x train
```

```
X_train.dtypes
```

```
Out[977]: gps_height          int64
population          int64
basin               category
month_recorded      category
region              category
public_meeting      int32
permit              int32
extraction_type_class category
management_group     category
payment              category
quality_group        category
quantity             category
source               category
waterpoint_type      category
construction_age     int32
dtype: object
```

We already know that there are no missing values, but we can check to confirm

```
In [978... X_train.isnull().sum()
```

```
Out[978]:  gps_height      0
  population      0
  basin           0
  month_recorded  0
  region          0
  public_meeting  0
  permit          0
  extraction_type_class  0
  management_group  0
  payment         0
  quality_group   0
  quantity        0
  source          0
  waterpoint_type  0
  construction_age  0
  dtype: int64
```

```
In [979... X_test.isnull().sum()
```

```
Out[979]:  gps_height      0
  population      0
  basin           0
  month_recorded  0
  region          0
  public_meeting  0
  permit          0
  extraction_type_class  0
  management_group  0
  payment         0
  quality_group   0
  quantity        0
  source          0
  waterpoint_type  0
  construction_age  0
  dtype: int64
```

Engineering outliers in numerical variables

We have seen that the `population` and `construction_age` , columns contain outliers

```
In [980... upper_thresholds = {
    'population': 560.0,
    'construction_age': 52.0,
}

for df3 in [X_train, X_test]:
    for column, top in upper_thresholds.items():
        df3[column] = df3[column].clip(upper=top)
```

```
In [981... print(round(X_train[clean_numeric_cols].describe()),2)
```

	gps_height	population	construction_age
count	47520.0	47520.0	47520.0
mean	669.0	184.0	27.0
std	693.0	153.0	10.0
min	-63.0	1.0	11.0
25%	0.0	100.0	20.0
50%	370.0	150.0	28.0
75%	1320.0	213.0	28.0
max	2770.0	560.0	52.0 2

```
In [982... print(round(X_test[clean_numeric_cols].describe()),2)
```

	gps_height	population	construction_age
count	11880.0	11880.0	11880.0
mean	667.0	185.0	27.0
std	694.0	152.0	10.0
min	-90.0	1.0	11.0
25%	0.0	100.0	20.0
50%	365.0	150.0	28.0
75%	1316.0	220.0	28.0
max	2623.0	560.0	52.0

We can see that the outliers in `population` and `construction_age` have been removed

Encode categorical variables

```
In [983... initial_object_cols
```

```
Out[983]: ['basin',
'month_recorded',
'region',
'extraction_type_class',
'management_group',
'payment',
'quality_group',
'quantity',
'source',
'waterpoint_type']
```

```
In [984... X_train[initial_object_cols].head()
```

```
Out[984]:
```

	basin	month_recorded	region	extraction_type_class	management_group	payment	quality_group
id							
454	Internal	2	Manyara	gravity	user-group	pay per bucket	good
510	Internal	3	Dodoma	handpump	user-group	never pay	good
14146	Lake Rukwa	7	Mbeya	other	user-group	never pay	good
47410	Rufiji	4	Mbeya	gravity	user-group	pay monthly	good
1288	Wami / Ruvu	4	Morogoro	other	user-group	pay when scheme fails	salty

```
In [985... X_train.head()
```

```
Out[985]:
```

	gps_height	population	basin	month_recorded	region	public_meeting	permit	extraction_type_cls
id								
454	2092	160.0	Internal	2	Manyara	1	1	gra
510	0	150.0	Internal	3	Dodoma	1	1	handpu
14146	0	150.0	Lake Rukwa	7	Mbeya	1	0	otl
47410	0	150.0	Rufiji	4	Mbeya	1	1	gra
1288	1023	120.0	Wami / Ruvu	4	Morogoro	1	1	otl


```
In [986... # using ohe Recommended

import category_encoders as ce

ohe = ce.one_hot.OneHotEncoder(cols=['basin','month_recorded','region', 'extraction_type',
                                     'quantity', 'source', 'waterpoint_type'], handle_missing="v")

ohe.fit(X_train)

X_train = ohe.transform(X_train)
X_test = ohe.transform(X_test)
```

Encode Status_group variable- Target

```
In [987... # import category_encoders as ce

# binary_encoder = ce.BinaryEncoder(cols=['RainToday'])

# X_train = binary_encoder.fit_transform(X_train)

# X_test = binary_encoder.transform(X_test)
```

```
In [988... X_train.head()
```

```
Out[988]:
```

	gps_height	population	basin_1	basin_2	basin_3	basin_4	basin_5	basin_6	basin_7	basin_8	...	s
id												
454	2092	160.0	1	0	0	0	0	0	0	0	...	
510	0	150.0	1	0	0	0	0	0	0	0	...	
14146	0	150.0	0	1	0	0	0	0	0	0	...	
47410	0	150.0	0	0	1	0	0	0	0	0	...	
1288	1023	120.0	0	0	0	1	0	0	0	0	...	

5 rows × 94 columns

```
In [989... X_test.head()
```

```
Out[989]:
```

	gps_height	population	basin_1	basin_2	basin_3	basin_4	basin_5	basin_6	basin_7	basin_8	...	s
id												
37098	0	150.0	0	0	0	0	0	1	0	0	...	
14530	0	150.0	0	0	0	0	0	1	0	0	...	
62607	1675	148.0	1	0	0	0	0	0	0	0	...	
46053	0	150.0	0	1	0	0	0	0	0	0	...	
47083	1109	235.0	1	0	0	0	0	0	0	0	...	

5 rows × 94 columns

We now have training and testing set ready for model building. Before that, we should map all the feature variables onto the same scale. It is called **feature scaling** . I will do it as follows.

In [990... `from sklearn.preprocessing import MinMaxScaler`

```
scaler = MinMaxScaler()

scaler.fit(X_train)
```

Out[990]: `MinMaxScaler`
`MinMaxScaler()`

In [991... `X_train = pd.DataFrame(`
 `scaler.transform(X_train),`
 `columns=X_train.columns`
 `)`

In [992... `X_test = pd.DataFrame(`
 `scaler.transform(X_test),`
 `columns=X_test.columns`
 `)`

In [993... `X_train.head()`

Out[993]:

	gps_height	population	basin_1	basin_2	basin_3	basin_4	basin_5	basin_6	basin_7	basin_8	...	source
0	0.760678	0.284436	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
1	0.022238	0.266547	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
2	0.022238	0.266547	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
3	0.022238	0.266547	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	...	
4	0.383339	0.212880	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	

5 rows × 94 columns

In [994... `X_train.isnull().sum()`

Out[994]:

gps_height	0
population	0
basin_1	0
basin_2	0
basin_3	0
...	..
waterpoint_type_4	0
waterpoint_type_5	0
waterpoint_type_6	0
waterpoint_type_7	0
construction_age	0

Length: 94, dtype: int64

In [995... `X_test.isnull().sum()`

Out[995]:

gps_height	0
population	0
basin_1	0
basin_2	0
basin_3	0
...	..
waterpoint_type_4	0
waterpoint_type_5	0
waterpoint_type_6	0
waterpoint_type_7	0
construction_age	0

Length: 94, dtype: int64

We now have `X_train` dataset ready to be fed into the Logistic Regression classifier. I will do it as follows.

Model Training

Model 1: Logistic Regression Iteration 1

```
In [996... # instantiate the model
logreg = LogisticRegression(solver='sag', max_iter=400, multi_class='auto', random_state

# fit the model
logreg.fit(X_train, y_train)
```

```
Out[996]: □ LogisticRegression
LogisticRegression(max_iter=400, random_state=42, solver='sag')
```

```
In [997... y_pred_test = logreg.predict(X_test)

y_pred_test
```

```
Out[997]: array([2, 0, 0, ..., 2, 0, 0], dtype=int64)
```

```
In [998... # Check accuracy score

from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'. format(accuracy_score(y_test, y_pred_test)))

Model accuracy score: 0.7297
```

Here, **y_test** are the true class labels and **y_pred_test** are the predicted class labels in the test-set.

Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
In [999... y_pred_train = logreg.predict(X_train)

y_pred_train
```

```
Out[999]: array([0, 0, 2, ..., 2, 0, 2], dtype=int64)
```

```
In [100... print('Training-set accuracy score: {0:0.4f}'. format(accuracy_score(y_train, y_pred_tra

Training-set accuracy score: 0.7322
```

Check for overfitting or underfitting

```
In [100... # print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))
```

Training set score: 0.7322
Test set score: 0.7297

The training-set accuracy score is 0.7311 while the test-set accuracy is 0.7297. These two values are quite comparable. So, there is no question of overfitting.

In Logistic Regression, we use default value of $C = 1$. It provides good performance with approximately 73% accuracy on both the training and the test set. But the model performance on both the training and test set are very comparable. It is likely the case of underfitting.

I will increase C and fit a more flexible model.

```
In [100... # fit the Logsitic Regression model with C=100
```

```
# instantiate the model
logreg100 = LogisticRegression(C=100, solver='sag', random_state=42)

# fit the model
logreg100.fit(X_train, y_train)
```

```
c:\Users\rosew\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
```

```
Out[1002]: □ LogisticRegression
LogisticRegression(C=100, random_state=42, solver='sag')
```

```
In [100... # print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(logreg100.score(X_train, y_train)))
print('Test set score: {:.4f}'.format(logreg100.score(X_test, y_test)))
```

```
Training set score: 0.7304
Test set score: 0.7294
```

We can see that, $C=100$ results in slightly test set accuracy and also a slightly reduced training set accuracy. So, we can conclude that a more complex model does not perform better.

Now, I will investigate, what happens if we use more regularized model than the default value of $C=1$, by setting $C=0.01$.

```
In [100... # fit the Logsitic Regression model with C=001
```

```
# instantiate the model
logreg001 = LogisticRegression(C=0.01, solver='sag', random_state=42)

# fit the model
logreg001.fit(X_train, y_train)
```

```
Out[1004]: □ LogisticRegression
LogisticRegression(C=0.01, random_state=42, solver='sag')
```

```
In [100... # print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(logreg001.score(X_train, y_train)))
```

```
print('Test set score: {:.4f}'.format(logreg001.score(X_test, y_test)))
```

Training set score: 0.7303

Test set score: 0.7307

The training and test set accuracy scores are also lower than the original, so we will continue from the one that used default $c = 1$

This will be our baseline model. Now let's try to change up some things and see what that does to our model

```
In [100... from sklearn.metrics import confusion_matrix

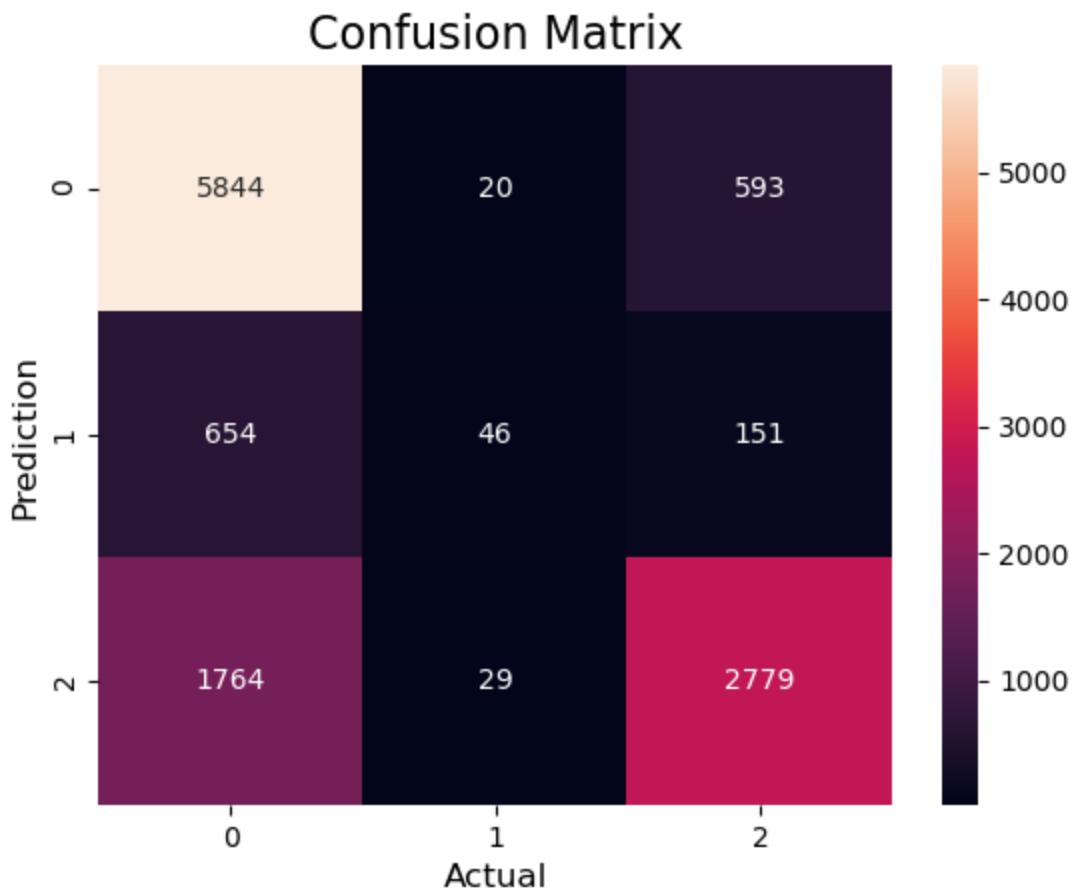
cm_logreg1 = confusion_matrix(y_test, y_pred_test)
print('Confusion matrix\n\n', cm_logreg1)

sns.heatmap(cm_logreg1, annot=True, fmt='d')
plt.ylabel('Prediction', fontsize=12)
plt.xlabel('Actual', fontsize=12)
plt.title('Confusion Matrix', fontsize=16)
# plt.show();
```

Confusion matrix

```
[[5844  20  593]
 [ 654  46  151]
 [1764  29 2779]]
Text(0.5, 1.0, 'Confusion Matrix')
```

Out[1006]:



Remember that:

0: Functional:

- 1: functional needs repair
- 2: non-functional

In [100...

```
from sklearn.metrics import classification_report  
  
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.71	0.91	0.79	6457
1	0.48	0.05	0.10	851
2	0.79	0.61	0.69	4572
accuracy			0.73	11880
macro avg	0.66	0.52	0.53	11880
weighted avg	0.72	0.73	0.70	11880

weighted : Calculate precision and recalls metrics for each label, and find their average weighted by support (the number of true instances for each label).

This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

In [100...

```
from sklearn.metrics import precision_score, recall_score, balanced_accuracy_score, f1_score  
  
precision = precision_score(y_test, y_pred_test, average='weighted')  
recall = recall_score(y_test, y_pred_test, average = "weighted")  
f1_score = f1_score(y_test, y_pred_test, average="weighted")  
  
accuracy_score = accuracy_score(y_test, y_pred_test)  
  
print("The weighted precision is: ",precision)  
print("The weighted recall is: ",recall)  
print("The accuracy score is: ", accuracy_score)  
print("The weighted f1score is: ", f1_score)  
print(" ")  
  
print("Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced")
```

```
The weighted precision is:  0.7227097828639878  
The weighted recall is:   0.7297138047138048  
The accuracy score is:   0.7297138047138048  
The weighted f1score is:  0.7027969482071396
```

Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced

Precision: 72%

Recall: 73%

The weighted average precision and recall are preferred as evaluation metrics.

It is appropriate to use weighted averaging. This approach takes into account the balance of classes. You weigh each class based on its representation in the dataset. Then, you compute precision and recall as a weighted average of the precision and recall in individual classes.

Simply put, it would work like macro-averaging, but instead of dividing precision and recall by the number of classes, you give each class a fair representation based on the proportion it takes in the dataset.

This approach is useful if you have an imbalanced dataset, like we have here but want to assign larger importance to classes with more examples.

Model 2- Logistic Regression Iteration Two

Our target variable is highly imbalanced. We will create a copy and use SMOTE to balance it then create a second iteration

```
In [100... Xy_init.status_group.value_counts()
```

```
Out[1009]: 0      32259
           2      22824
           1       4317
           Name: status_group, dtype: int64
```

```
In [101... Xy_init_smote = Xy_init.copy() # assign to protect original one

X_smote = Xy_init_smote.drop(['status_group'], axis=1) # assign X variables

y_smote = Xy_init_smote['status_group'] # Assign y variable

from sklearn.model_selection import train_test_split

X_train2, X_test2, y_train2, y_test2 = train_test_split(X_smote, y_smote, test_size = 0.

# Remove outliers for numerical data

upper_thresholds = {
    'population': 560.0,
    'construction_age': 52.0,
}

for df3 in [X_train2, X_test2]:
    for column, top in upper_thresholds.items():
        df3[column] = df3[column].clip(upper=top)

# Obe Hot encode categorical data

import category_encoders as ce

ohe = ce.one_hot.OneHotEncoder(cols=['basin', 'month_recorded', 'region', 'extraction_type',
                                     'quantity', 'source', 'waterpoint_type'], handle_missing="v

ohe.fit(X_train2)

X_train2 = ohe.transform(X_train2)
X_test2 = ohe.transform(X_test2)

# Scale the data

from sklearn.preprocessing import MinMaxScaler

scaler2 = MinMaxScaler()

scaler2.fit(X_train2)

X_train2 = pd.DataFrame(
    scaler.transform(X_train2),
    columns=X_train2.columns
)

X_test2 = pd.DataFrame(
    scaler.transform(X_test2),
```

```

        columns=X_test2.columns
    )

    # SMOTE
    from imblearn.over_sampling import SMOTE

    smt = SMOTE(sampling_strategy = 'auto', n_jobs = -1)

    X_smote_sampled, y_smote_sampled = smt.fit_resample(X_train2, y_train2)

    print(y_smote.value_counts())

    y_smote_sampled = pd.Series(y_smote_sampled) # converting from array to np.series to see
    print(y_smote_sampled.value_counts())


    # # Instantiate and Fit Logistic Regression Model
    # # instantiate the model
    logreg2 = LogisticRegression(solver='sag', multi_class='auto', random_state=42)

    # # fit the model
    logreg2.fit(X_smote_sampled, y_smote_sampled)

```

```

0      32259
2      22824
1       4317
Name: status_group, dtype: int64
2      25802
1      25802
0      25802
Name: status_group, dtype: int64

```

```

c:\Users\rosew\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

```

```

Out[1010]: □ LogisticRegression
LogisticRegression(random_state=42, solver='sag')

```

```

In [101... y_pred_test2 = logreg2.predict(X_test2)

y_pred_test2

```

```

Out[1011]: array([2, 0, 0, ..., 2, 1, 0], dtype=int64)

```

```

In [101... from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test2, y_pred_test2)))

Model accuracy score: 0.6215

```

Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```

In [101... y_pred_train2 = logreg2.predict(X_train2)

y_pred_train2

```



```
Out[1013]: array([0, 0, 2, ..., 2, 1, 2], dtype=int64)
```

```
In [101... print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train2, y_pred_tr
```

Training-set accuracy score: 0.6273

Check for overfitting or underfitting

```
In [101... # print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(logreg2.score(X_train2, y_train2)))
```

```
print('Test set score: {:.4f}'.format(logreg2.score(X_test2, y_test2)))
```

Training set score: 0.6273

Test set score: 0.6215

The training and test accuracy score are much lower than the first model (62%), meaning that it is less accurate. Using SMOTE does not improve our model

```
In [101... from sklearn.metrics import confusion_matrix
```


```
cm_logreg3 = confusion_matrix(y_test2, y_pred_test2)
print('Confusion matrix\n\n', cm_logreg3)
```

```
sns.heatmap(cm_logreg3, annot=True, fmt='d')
plt.ylabel('Prediction', fontsize=12)
plt.xlabel('Actual', fontsize=12)
plt.title('Confusion Matrix', fontsize=16)
plt.show();
```

Confusion matrix

```
[[4040 1724  693]
 [ 209  538  104]
 [ 953  814 2805]]
```

Confusion Matrix



```
In [101]: from sklearn.metrics import classification_report
print(classification_report(y_test2, y_pred_test2))
```

	precision	recall	f1-score	support
0	0.78	0.63	0.69	6457
1	0.17	0.63	0.27	851
2	0.78	0.61	0.69	4572
accuracy			0.62	11880
macro avg	0.58	0.62	0.55	11880
weighted avg	0.73	0.62	0.66	11880

weighted : Calculate precision and recalls metrics for each label, and find their average weighted by support (the number of true instances for each label).

This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

```
In [101]: from sklearn.metrics import precision_score, recall_score, balanced_accuracy_score, f1_score

precision2 = precision_score(y_test, y_pred_test2, average='weighted')
recall2 = recall_score(y_test, y_pred_test2, average = "weighted")
f1_score2 = f1_score(y_test, y_pred_test2, average="weighted")

accuracy_score2 = accuracy_score(y_test, y_pred_test2)

print("The weighted precision is: ",precision2)
print("The weighted recall is: ",recall2)
print("The accuracy score is: ", accuracy_score2)
print("The weighted f1score is: ", f1_score2)
print(" ")
```

```
print("Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced")
```

The weighted precision is: 0.7343331320769583

The weighted recall is: 0.6214646464646465

The accuracy score is: 0.6214646464646465

The weighted f1score is: 0.6604305711351062

Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced

Precision: 74%

Recall: 62%

Let's now try using the standard scaler instead on the MinMax Scaler on the original data to see if this creates a better model

Model 3: Linear Regression Iteration 3

```
In [101]: Xy_init_std_scaler = Xy_init.copy() # assign to protect original one

X_std_scaler = Xy_init_std_scaler.drop(['status_group'], axis=1) # assign X variables
y_std_scaler = Xy_init_std_scaler['status_group'] # Assign y variable

from sklearn.model_selection import train_test_split

X_train3, X_test3, y_train3, y_test3 = train_test_split(X_std_scaler, y_std_scaler, test_size=0.2,
                                                        # Remove outliers for numerical data

                                                        upper_thresholds = {
                                                            'population': 560.0,
                                                            'construction_age': 52.0,
                                                        })

for df3 in [X_train3, X_test3]:
    for column, top in upper_thresholds.items():
        df3[column] = df3[column].clip(upper=top)

# One Hot encode categorical data

import category_encoders as ce

ohe = ce.one_hot.OneHotEncoder(cols=['basin', 'month_recorded', 'region', 'extraction_type',
                                     'quantity', 'source', 'waterpoint_type'], handle_missing='value')

ohe.fit(X_train3)

X_train3 = ohe.transform(X_train3)
X_test3 = ohe.transform(X_test3)

# Scale the data

from sklearn.preprocessing import StandardScaler

scaler3 = StandardScaler()

scaler3.fit(X_train3)

X_train3 = pd.DataFrame(
    scaler3.transform(X_train3),
```

```

        columns=X_train3.columns
    )

X_test3 = pd.DataFrame(
    scaler3.transform(X_test3),
    columns=X_test3.columns
)

## Instantiate and Fit Logistic Regression Model
## instantiate the model
logreg3 = LogisticRegression(solver='sag', multi_class='auto', random_state=42)

## fit the model
logreg3.fit(X_train3, y_train3)

```

```

Out[1019]: □ LogisticRegression
LogisticRegression(random_state=42, solver='sag')

```

```

In [102... y_pred_test3 = logreg3.predict(X_test3)

y_pred_test3

```

```

Out[1020]: array([2, 0, 0, ..., 2, 0, 0], dtype=int64)

```

```

In [102... from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test3, y_pred_test3)))

Model accuracy score: 0.7295

```

Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```

In [102... ### Compare the train-set and test-set accuracy

# Now, I will compare the train-set and test-set accuracy to check for overfitting.
y_pred_train3 = logreg3.predict(X_train3)

y_pred_train3

print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train3, y_pred_train3)))
# Check for overfitting or underfitting
# print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg3.score(X_train3, y_train3)))

print('Test set score: {:.4f}'.format(logreg3.score(X_test3, y_test3)))

Training-set accuracy score: 0.7323
Training set score: 0.7323
Test set score: 0.7295

```

The training and test accuracy score are similar to the first model at 0.73 and 0.729

```

In [102... from sklearn.metrics import confusion_matrix

```

```

cm_logreg4 = confusion_matrix(y_test3, y_pred_test3)
print('Confusion matrix\n\n', cm_logreg4)

sns.heatmap(cm_logreg4, annot=True, fmt='d')
plt.ylabel('Prediction', fontsize=12)
plt.xlabel('Actual', fontsize=12)
plt.title('Confusion Matrix', fontsize=16)
plt.show();


```

Confusion matrix

```

[[5842  23  592]
 [ 654  46  151]
 [1764  29 2779]]

```



```

In [102... from sklearn.metrics import classification_report

print(classification_report(y_test3, y_pred_test3))

```

	precision	recall	f1-score	support
0	0.71	0.90	0.79	6457
1	0.47	0.05	0.10	851
2	0.79	0.61	0.69	4572
accuracy			0.73	11880
macro avg	0.66	0.52	0.53	11880
weighted avg	0.72	0.73	0.70	11880

weighted : Calculate precision and recalls metrics for each label, and find their average weighted by support (the number of true instances for each label).

This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

```

from sklearn.metrics import precision_score, recall_score, balanced_accuracy_score, f1_s

```

```
In [102... precision3 = precision_score(y_test, y_pred_test3, average='weighted')
recall3 = recall_score(y_test, y_pred_test3, average = "weighted")
f1_score3 = f1_score(y_test, y_pred_test3, average="weighted")

accuracy_score3 = accuracy_score(y_test, y_pred_test3)

print("The weighted precision is: ",precision3)
print("The weighted recall is: ",recall3)
print("The accuracy score is: ", accuracy_score3)
print("The weighted f1score is: ", f1_score3)
print(" ")

print("Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced")

The weighted precision is: 0.7216956617165646
The weighted recall is: 0.7295454545454545
The accuracy score is: 0.7295454545454545
The weighted f1score is: 0.702718498868645

Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced

Precision: 72%
Recall: 73%
```

Decision Tree

Model 4: Decision Tree- Model Iteration 1

```
In [102... from sklearn import tree
from sklearn.tree import DecisionTreeClassifier

In [102... dt = DecisionTreeClassifier(criterion='gini', splitter='best', random_state=42)

In [102... dt.fit(X_train, y_train)

Out[1028]: □ DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

Compare the train-set and test-set accuracy

Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
In [102... print('Accuracy score of train data :{}'.format(dt.score(X_train,y_train)))
print('Accuracy score of test data:{}'.format(dt.score(X_test,y_test)))

Accuracy score of train data :0.937962962962963
Accuracy score of test data:0.7570707070707071
```

This high variance in accuracy scores indicates that we are overfitting. Let's change the depth and use entropy and see the difference it brings

```
In [103... y_pred_test4 = dt.predict(X_test)

y_pred_test4
```

Out[1030]: array([2, 0, 0, ..., 2, 2, 1], dtype=int64)

```
In [103... from sklearn.metrics import confusion_matrix


cm_logreg2 = confusion_matrix(y_test, y_pred_test4)
print('Confusion matrix\n\n', cm_logreg2)

sns.heatmap(cm_logreg2, annot=True, fmt='d')

plt.ylabel('Prediction', fontsize=12)
plt.xlabel('Actual', fontsize=12)
plt.title('Confusion Matrix', fontsize=16)
plt.show();
```

Confusion matrix

```
[[5359  316  782]
 [ 428  280  143]
 [1066  151 3355]]
```



```
In [103... from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred_test4))
```

	precision	recall	f1-score	support
0	0.78	0.83	0.81	6457
1	0.37	0.33	0.35	851
2	0.78	0.73	0.76	4572
accuracy			0.76	11880
macro avg	0.65	0.63	0.64	11880
weighted avg	0.75	0.76	0.75	11880

`weighted` : Calculate precision and recalls metrics for each label, and find their average weighted by support (the number of true instances for each label).

This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

```
In [103... from sklearn.metrics import precision_score, recall_score, balanced_accuracy_score, f1_score

precision4 = precision_score(y_test, y_pred_test4, average='weighted')
recall4 = recall_score(y_test, y_pred_test4, average = "weighted")
f1_score4 = f1_score(y_test, y_pred_test4, average="weighted")

accuracy_score4 = accuracy_score(y_test, y_pred_test4)

print("The weighted precision is: ",precision4)
print("The weighted recall is: ",recall4)
print("The accuracy score is: ", accuracy_score4)
print("The weighted f1score is: ", f1_score4)
print(" ")

print("Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced")

The weighted precision is:  0.7535526749600432
The weighted recall is:  0.7570707070707071
The accuracy score is:  0.7570707070707071
The weighted f1score is:  0.7544993679253735
```

Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced

Model 5: Decision Tree- Model Iteration Two: Hyperparameter Tuning for the Decision Tree

```
In [103... dt2= DecisionTreeClassifier(criterion='entropy', splitter='best', max_depth= 14, random_state=42)
dt2.fit(X_train, y_train)

### Compare the train-set and test-set accuracy
# Now, I will compare the train-set and test-set accuracy to check for overfitting.

print('Accuracy score of train data :{}'.format(dt2.score(X_train,y_train)))
print('Accuracy score of test data:{}'.format(dt2.score(X_test,y_test)))

Accuracy score of train data :0.7965698653198653
Accuracy score of test data:0.7576599326599327
```

```
In [103... y_pred_test5 = dt2.predict(X_test)

y_pred_test5

Out[1035]: array([2, 0, 0, ..., 2, 0, 0], dtype=int64)
```

Accuracy = 75.8% .

This has improved both accuracy scores and we are not overfitting either. This could be perfect for the final model

Let's create a confusion matrix for this and obtain the other classification metrics

```
In [103... from sklearn.metrics import confusion_matrix

cm_logreg5 = confusion_matrix(y_test, y_pred_test5)
```




```
print('Confusion matrix\\n\\n', cm_logreg5)

sns.heatmap(cm_logreg5, annot=True,fmt='d')

plt.ylabel('Prediction',fontSize=12)
plt.xlabel('Actual',fontSize=12)
plt.title('Confusion Matrix',fontSize=16)
plt.show();
```

Confusion matrix

```
[[5789   94  574]
 [ 533  177  141]
 [1458   79 3035]]
```



Remember that:

0: Functional:

1: functional needs repair

2: non-functional

In [103...

```
from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred_test5))
```

	precision	recall	f1-score	support
0	0.74	0.90	0.81	6457
1	0.51	0.21	0.29	851
2	0.81	0.66	0.73	4572
accuracy			0.76	11880
macro avg	0.69	0.59	0.61	11880
weighted avg	0.75	0.76	0.74	11880

weighted : Calculate precision and recalls metrics for each label, and find their average weighted by

support (the number of true instances for each label).

This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

In [103...

```
from sklearn.metrics import precision_score, recall_score, balanced_accuracy_score, f1_score

precision5 = precision_score(y_test, y_pred_test5, average='weighted')
recall5 = recall_score(y_test, y_pred_test5, average = "weighted")
f1_score5 = f1_score(y_test, y_pred_test5, average="weighted")

accuracy_score5 = accuracy_score(y_test, y_pred_test5)

print("The weighted precision is: ",precision5)
print("The weighted recall is: ",recall5)
print("The accuracy score is: ", accuracy_score5)
print("The weighted f1score is: ", f1_score5)
print(" ")

print("Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced")

The weighted precision is:  0.7521218200489409
The weighted recall is:  0.7576599326599327
The accuracy score is:  0.7576599326599327
The weighted f1score is:  0.7438267978791073
```

Weighted precision, f1 score, and recall are ideal measures because the target labels are imbalanced

Precision = 75%

Recall = 76%

Conclusion

1. The relationship between the following variables and the `status_group` (functional, non-functional, functional but needs repair):

Wells likely to be non-functional or needing repair:

- `payment` : Wells where no payments are made
- `source` : Wells with a shallow well as the water source
- `management_group` : Wells managed by the user group
- `extraction_type` : Wells with gravity as the extraction type
- `permit` : Wells that are permitted
- `public_meeting` : Wells where a public meeting was held

2. Modeling

The best model is the final one- Decision Tree Classifier with a maximum depth of 14 and entropy criterion
Its metrics are:

- Accuracy: 75.8%
- Precision = 75%
- Recall- 76%

Recommendations

- Consider engineering the ternary classification problem into a binary classification problem and see if this improves the model parameters
- More feature engineering to identify other features that could be useful for the model.
- Try other models like the Random Forest Classifier, KNN, and AdaBoost to see if they have better metrics

In []: